

Asymmetric cryptography

Asymmetric (Block) Ciphers

Use key pairs

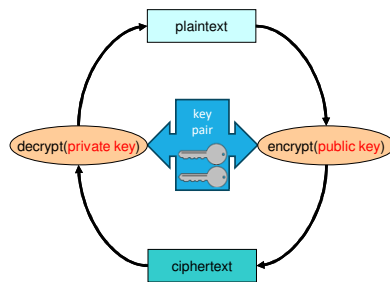
- One private key (personal, not transmittable)
- One public key, available to all

Allow

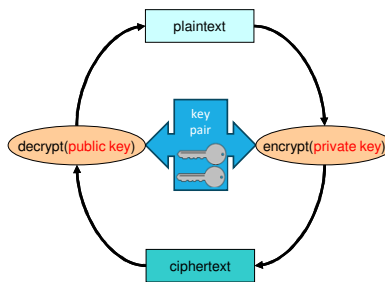
- Confidentiality without any previous exchange of secrets
- Authentication
 - Of contents (data integrity)
 - Of origin (source authentication, or digital signature)

Operations of an asymmetric cipher

Confidentiality



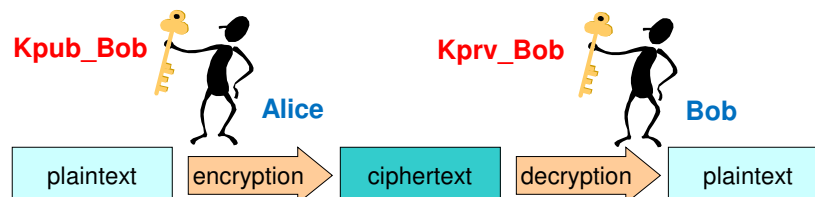
Authentication (signature)



Use cases: secure communication

Secure communication with a target (Bob)

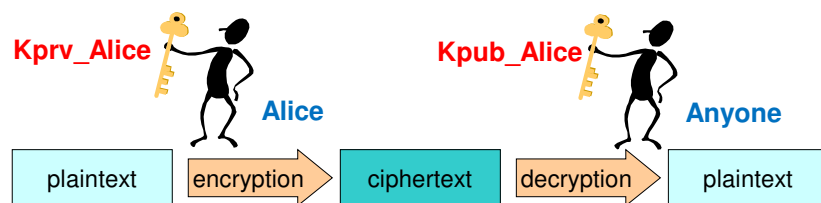
- Alice encrypts plaintext **P** with Bob's public key **Kpub_Bob**
Alice: $C = \{P\}_{K_{pub_Bob}}$
- Bob decrypts cyphertext **C** with his private key **Kprv_Bob**
Bob: $P' = \{C\}_{K_{prv_Bob}}$
- **P'** should be equal to **P** (requires checking)
- **Kpub_Bob** needs to be known by Alice



Use cases: signature

Data signature by Alice

- Alice encrypts plaintext **P** with her private key **K_{priv_Alice}**
Alice: $C = \{P\}_{K_{priv_Alice}}$
- Anyone can decrypt cyphertext **C** with Alice's public key **K_{pub_Alice}**
Anyone: $P' = \{C\}_{K_{pub_Bob}}$
- If $P' = P$, then **C** is Alice's signature of **P**
- **K_{pub_Alice}** needs to be known by signature verifiers



Asymmetric ciphers

Advantages

- They are a fundamental authentication mechanism
- They allow to explore features that are not possible with asymmetric ciphers

Disadvantages

- Performance
- Usually are very inefficient and memory consuming

Problems

- Trustworthy distribution of public keys
- Lifetime of key pairs

Asymmetric ciphers

Approaches: complex mathematic problems

- Discrete logarithms of large numbers
- Integer factorization of large numbers

Most common algorithms

- RSA
- ElGamal
- Elliptic curves (ECC)

Other techniques with asymmetric key pairs

- Diffie-Hellman (key agreement)

RSA (Rivest, Shamir, Adelman, 1978)

Keys

- Private: (d, n)
- Public: (e, n)

Public key encryption (confidentiality)

- $C = P^e \bmod n$
- $P = C^d \bmod n$

P, C are numbers

$0 \leq P, C < n$

Private key encryption (signature)

- $C = P^d \bmod n$
- $P = C^e \bmod n$

RSA (Rivest, Shamir, Adelman, 1978)

Computational complexity

- Discrete logarithm
- Integer factoring

coprime $\rightarrow \gcd(a, b) = 1$
 $\times \rightarrow$ multiplication
 $\text{mod} \rightarrow$ modulo operation
 $\equiv \rightarrow$ modular congruence

Key selection

- Large n (hundreds or thousands of bits)
- $n = p \times q$ with p and q being large (secret) prime numbers
- Chose an e co-prime with $(p-1) \times (q-1)$
- Compute d such that $e \times d \equiv 1 \pmod{(p-1) \times (q-1)}$
- Discard p and q
- The value of d cannot be computed out of e and n
 - Only from p and q

RSA example

$p = 5$ $q = 11$ (prime numbers)

- $n = p \times q = 55$
- $(p-1) \times (q-1) = 40$

$e = 3$ (public key = e, n)

- Coprime of 40

$d = 27$ (private key = d, n)

- $e \times d \equiv 1 \pmod{40} \rightarrow d \times e \text{ mod } 40 = 1, (27 \times 3) \text{ mod } 40 = 1$

For $P = 26$ (notice that $P, C \in [0, n-1]$)

- $C = P^e \text{ mod } n = 26^3 \text{ mod } 55 = 31$
- $P = C^d \text{ mod } n = 31^{27} \text{ mod } 55 = 26$

Hybrid encryption

Combines symmetric with asymmetric cryptography

- Use the best of both worlds, while avoiding problems
- Asymmetric cipher: Uses public keys (but it is slow)
- Symmetric cipher: Fast (but with weak key exchange methods)

Method:

- Obtain K_{pub} from the receiver
- Generate a random K_{sym}
- Calculate $C1 = E_{sym}(K_{sym}, P)$
- Calculate $C2 = E_{asym}(K_{pub}, K_{sym})$
- Send $C1 + C2$
 - $C1$ = Text encrypted with symmetric key
 - $C2$ = Symmetric key encrypted with the receiver public key
 - May also contain the IV

Randomization of asymmetric encryptions

Non-deterministic (unpredictable) result of asymmetric encryptions

- **N** encryptions of the same value, with the same key, should yield **N** different results
- **Goal:** prevent the trial & error discovery of encrypted values

Approaches

- Concatenation of value to encrypt with two values
 - A fixed one (for integrity control)
 - A random one (for randomization)

Randomization of asymmetric encryptions: OAEP (Optimal Asymmetric Encryption Padding)

IHash: digest over **Label**

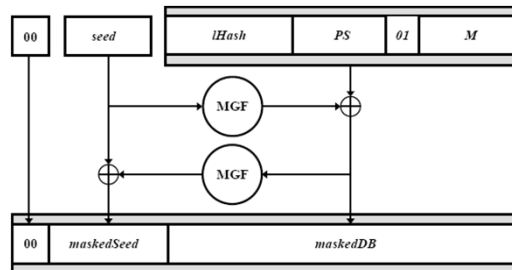
seed: random

PS: zeros

M: plaintext

MGF: Mask Generation Function

- Similar to Hash, but with variable size



Diffie-Hellman Key Agreement (1976)



q (large prime)
 α (primitive root mod q)



a = random

$$Y_a = \alpha^a \text{ mod } q$$

$$K_{ab} = Y_b^a \text{ mod } q$$



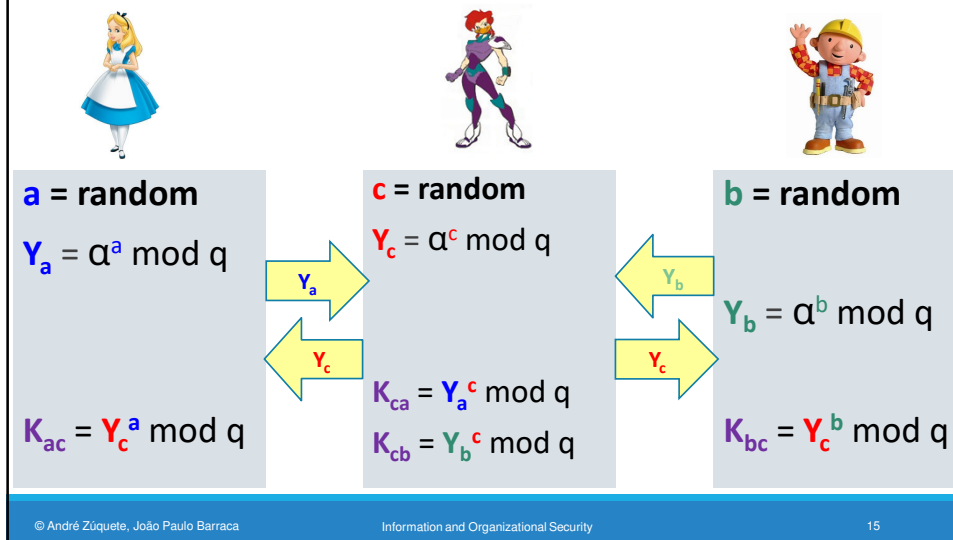
b = random

$$Y_b = \alpha^b \text{ mod } q$$

$$K_{ba} = Y_a^b \text{ mod } q$$

$$K_{ab} = K_{ba}$$

DH Key Agreement: MitM attack



Elliptic Curve Cryptography (ECC)

Elliptic curves are specific functions

- They have a generator (G)
- A private key K_{prv} is an integer with a maximum of bits allowed by the curve
- A public key K_{pub} is a point $(x, y) = K_{\text{prv}} \times G$
- Given K_{pub} , it should be hard to guess K_{prv}

Curves

- NIST curves (15)
 - P-192, P-224, P-256, P-384, P-521
 - B-163, B-233, B-283, B-409, B-571
 - K-163, K-233, K-283, K-409, K-571

Other curves

- Curve25519 (256 bits)
- Curve448 (448 bits)

ECDH: DH with ECC



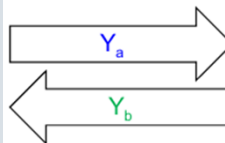
ECC curve $\rightarrow G$



a = random

$$Y_a = a G$$

$$K_{ab} = a Y_b$$



b = random

$$Y_b = b G$$

$$K_{ba} = b Y_a$$

$$K_{ab} = K_{ba}$$

ECC public key encryption

Combines hybrid encryption with ECDH

Method:

- Obtain K_{pub_recv} from the receiver
- Generate a random K_{prv_send} and the corresponding K_{pub_send}
- Calculate $K_{sym} = K_{prv_send} K_{pub_recv}$
- $C = E(P, K_{sym})$
- Send $C + K_{pub_send}$
- Receiver calculates $K_{sym} = K_{pub_send} K_{prv_recv}$
- $P = D(C, K_{sym})$

Digital signatures

Asymmetric (Block) Ciphers

Use key pairs

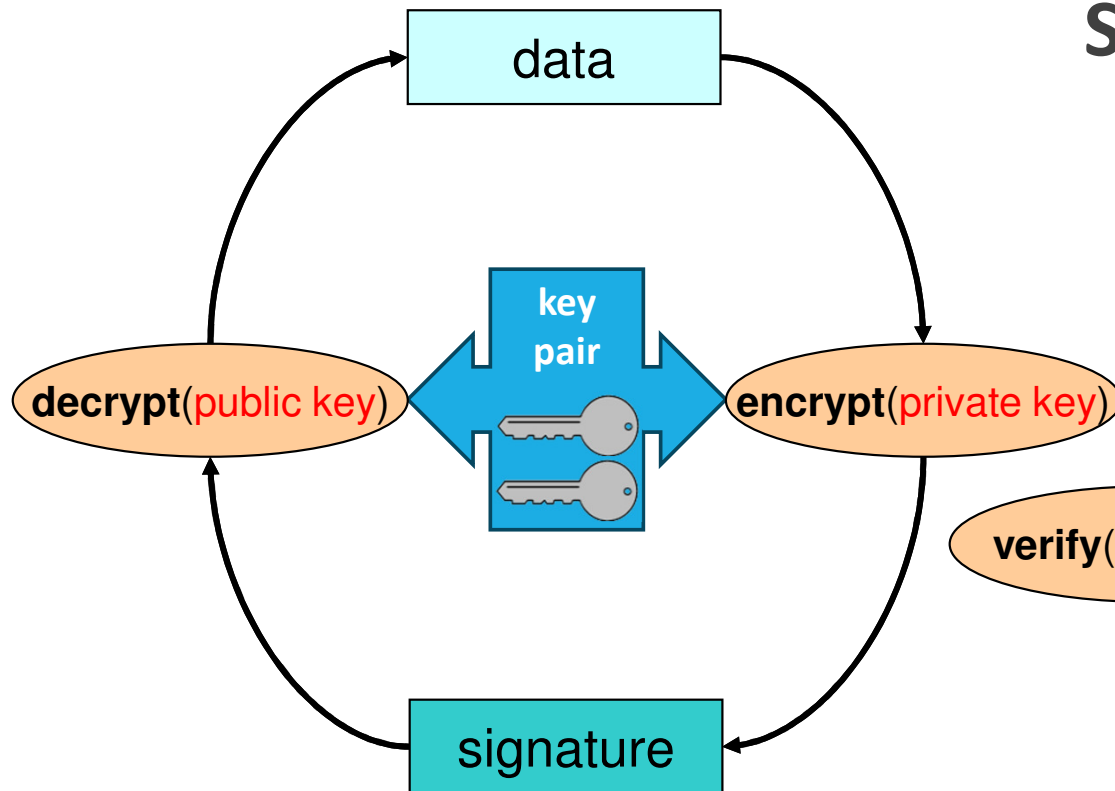
- One private key (personal, not transmittable)
- One public key, available to all

Allow

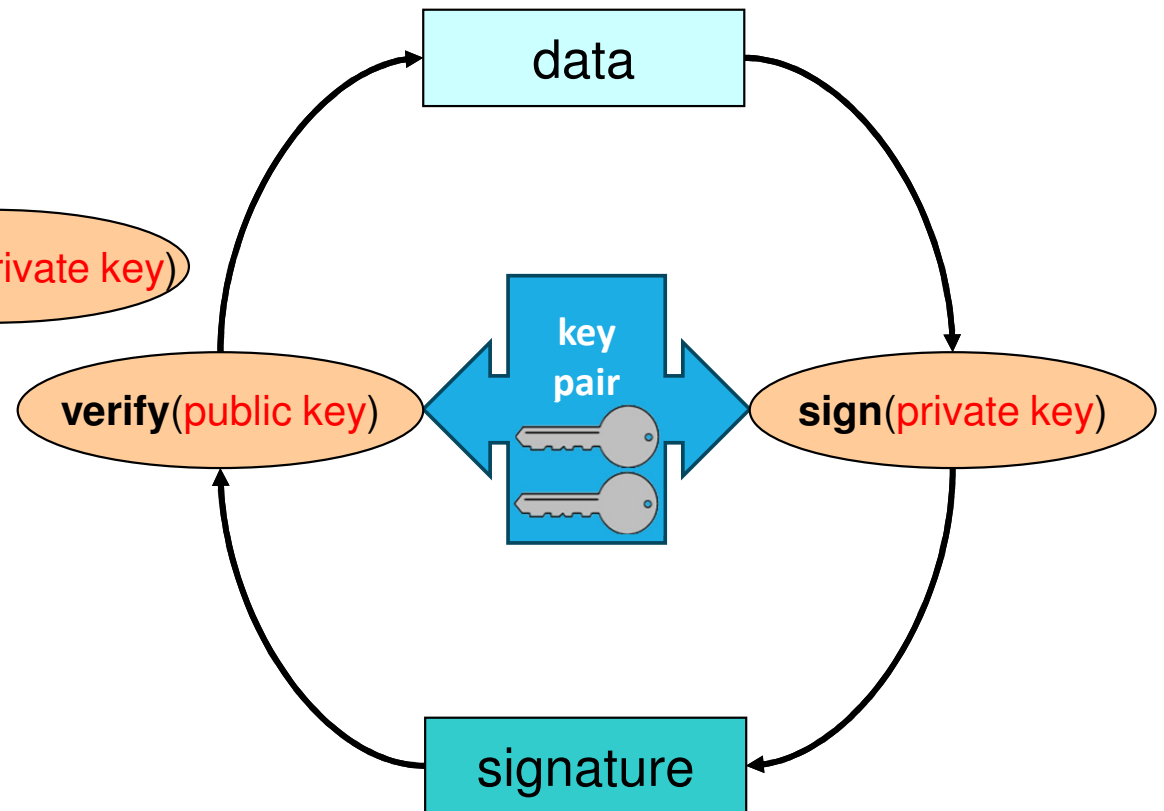
- Confidentiality without any previous exchange of secrets
- Authentication
 - Of contents (data integrity)
 - Of origin (source authentication, or digital signature)

Digital signatures

Encrypt / decrypt (RSA)



Sign / verify (ElGamal, EC)



Digital signatures

Authenticate the contents of a document

- Ensure its integrity (it was not changed)

Authenticate its author

- Ensure the identity of the creator/originator

Prevent repudiation of signatures

- Non-repudiation
- Genuine authors cannot deny authorship
 - Only the identified author could have generated a given signature

Digital Signatures

Approaches

- Asymmetric encryption/decryption or signature/verification
- Digest functions (only for performance)

Signing: $A_x(\text{doc}) = \text{info} + E(K_x^{-1}, \text{digest}(\text{doc} + \text{info}))$

$A_x(\text{doc}) = \text{info} + S(K_x^{-1}, \text{digest}(\text{doc} + \text{info}))$

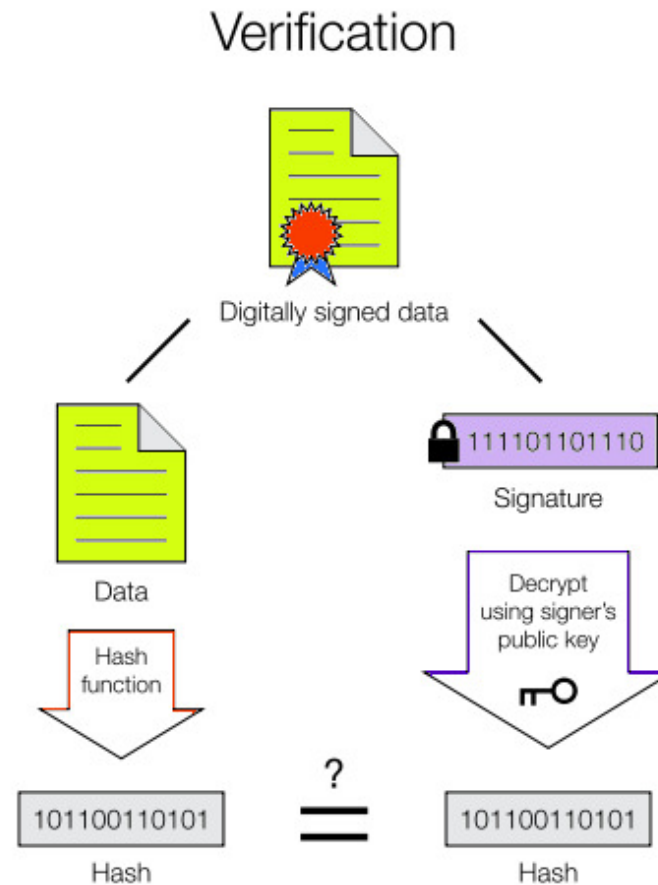
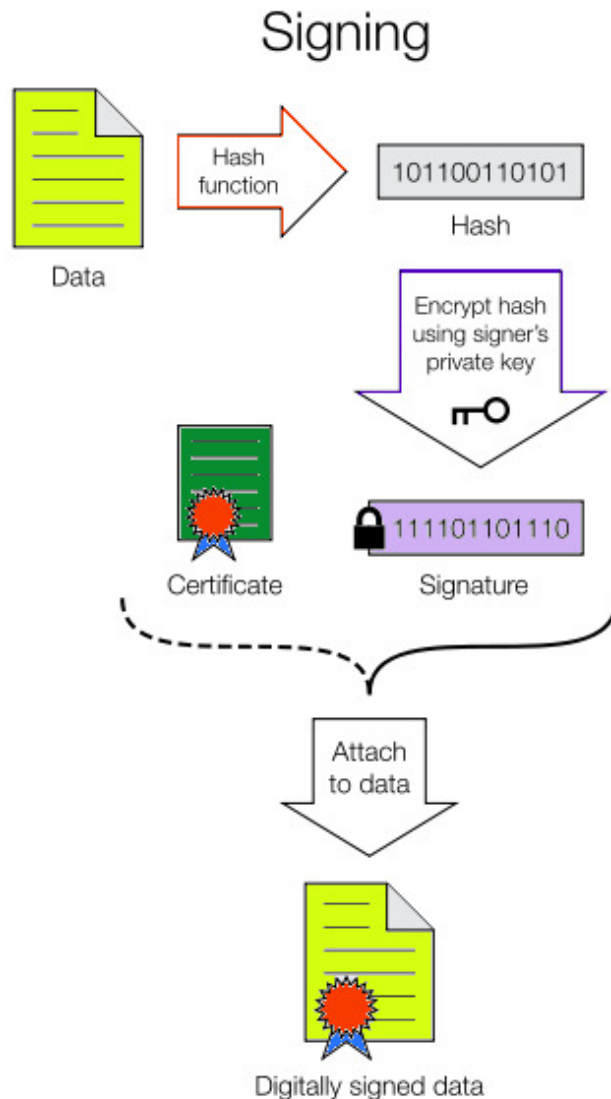
info = signing context, signer identity, K_x

Verification:

$D(K_x, A_x(\text{doc})) \equiv \text{digest}(\text{doc} + \text{info})$

$V(K_x, A_x(\text{doc}), \text{doc}, \text{info}) \rightarrow \text{True} / \text{False}$

Encryption / decryption signatures



If the hashes are equal, the signature is valid.

Digital signature on a mail:

Multipart content, signature w/ certificate

```
From - Fri Oct 02 15:37:14 2009
[...]  
Date: Fri, 02 Oct 2009 15:35:55 +0100  
From: =?ISO-8859-1?Q?Andr=E9_Z=FAquete?= <andre.zuquete@ua.pt>  
Reply-To: andre.zuquete@ua.pt  
Organization: IEETA / UA  
MIME-Version: 1.0  
To: =?ISO-8859-1?Q?Andr=E9_Z=FAquete?= <andre.zuquete@ua.pt>  
Subject: Teste  
Content-Type: multipart/signed; protocol="application/x-pkcs7-signature"; micalg=sha1; boundary="-----ms050405070101010502050101"
```

This is a cryptographically signed message in MIME format.

```
-----ms050405070101010502050101  
Content-Type: multipart/mixed;  
boundary="-----060802050708070409030504"
```

This is a multi-part message in MIME format.
-----060802050708070409030504
Content-Type: text/plain; charset=ISO-8859-1
Content-Transfer-Encoding: quoted-printable

Corpo do mail

```
-----060802050708070409030504--  
-----ms050405070101010502050101  
Content-Type: application/x-pkcs7-signature; name="smime.p7s"  
Content-Transfer-Encoding: base64  
Content-Disposition: attachment; filename="smime.p7s"  
Content-Description: S/MIME Cryptographic Signature
```

```
MIAGCSqGSIB3DQEHAQcAMIAQAQExCzAJBgUrDgMCGGUAMIAAGCSqGSIB3DQEHAQAooIIamTCC  
BUkwggSyoAMCAQICBAcnIaEwDQYJKoZIhvcNAQEFBQAwdTELMAkGA1UEBhMCVVMxGDAWBGNV  
[...]  
KoZIhvcNAQEBBQAEgYCOfs852BV77NVuww53vSx01XtI2JhC1CDlu+tcTPoMD1wq5dc5v40  
Tgsaw0N8dqgVLk8aC/CdGMbRBu+J1LKrcVZa+khnjtB66HhDRLrjmEGDNttrEjbbqvpd2Q02  
vxB3iPTIU+vCGXo47e6GyRydqTpbq0r49Zqmx+IJ6Z7iigAAAAA==  
-----ms050405070101010502050101--
```


Key derivation

Key derivation

Cipher algorithms require fixed dimension keys

- 56, 128, 256... bits

We may derive keys from multiple sources

- Shared secrets
- Passwords generated by humans
- PIN codes and small length secrets

Original source may have low entropy

- Reduces the difficulty of a brute force attack
- Although we must have some strong relation into a useful key

Sometimes we need multiple keys from the same material

- While not allowing to find the material (a password, another key) from the new key

Key derivation: purposes

Key reinforcement: increase the security of a password

- Usually defined by humans
- Making dictionary attacks impractical

Key expansion: increase the dimension of a key

- Expansion to a size that suits an algorithm
- Eventually derive other related keys for other algorithms (e.g. MAC)

Key derivation

Key derivation requires the existence of:

- A salt which makes the derivation unique
- A difficult problem
- A chosen level of complexity

Computational difficulty

- Transformation requires relevant computational resources

Memory difficulty

- Transformation requires relevant storage resources
- Limits attacks using dedicated hardware accelerators

Key derivation: PKBDF2

Password Based Key Derivation Function 2

Produces a key from a password, with a chosen difficulty

$K = \text{PBKDF2}(\text{PRF}, \text{Salt}, \text{rounds}, \text{dim}, \text{password})$

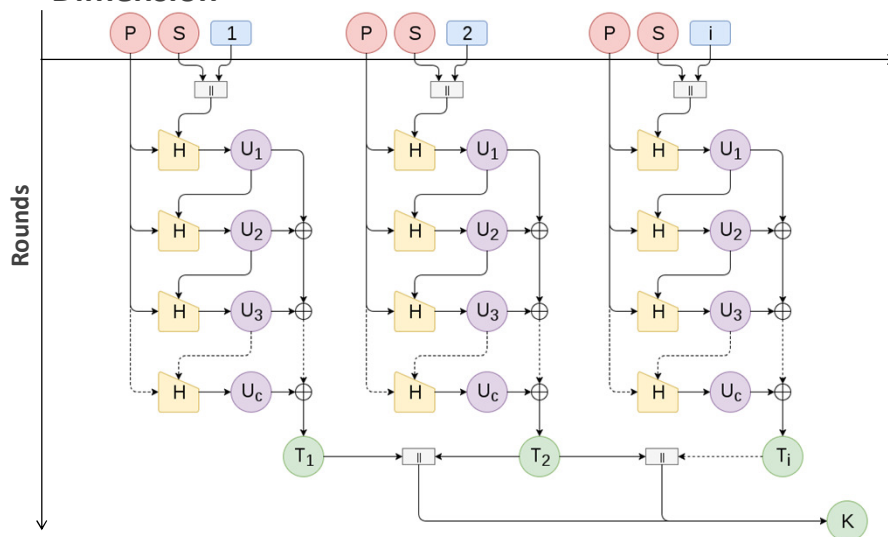
- PRF: Pseudo-Random-Function: a digest function
- Salt: a random value
- Rounds: the computational cost (tens or hundreds of thousands)
- Dim: the size of the result required

Operation: calculates $\text{ROUNDS} \times \text{DIM}$ operations from the PRF using the SALT and PASSWORD

- Larger number of rounds will increase the cost

Key Derivation: PBKDF2

Dimension



Key Derivation: scrypt

Produces a key with a chosen storage cost

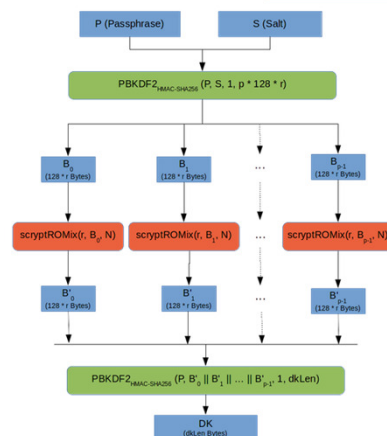
$K = \text{scrypt}(\text{password}, \text{salt}, n, p, \text{dim}, r, \text{hLen}, \text{Mflen})$

- Password: a secret
- Salt: a random value
- N: the cost parameter
- P: the parallelization parameter. $p \leq (2^{32} - 1) * \text{hLen} / \text{Mflen}$
- Dim: the size of the result
- R: the size of the blocks to use (default is 8)
- hLen: the size of the digest function (32 for SHA256)
- Mflen: bytes in the internal mix (default is $8 \times R$)

Key Derivation: scrypt

$\text{scrypt}(P, S, N, r, p, \text{dkLen})$

Parameters:
N (CPU/Memory Cost Parameter)
r (Block Size)
p (Parallelization Parameter)
dkLen (Output Length)



Management of Asymmetric keys

Problems to solve

Ensure proper and correct use of asymmetric key pairs

Privacy of private keys

- To ensure authenticity
- To prevent the repudiation of digital signatures

Correct distribution of public keys

- To ensure confidentiality
- To ensure the correct validation of digital signatures

Problems to solve

Temporal evolution of entity <-> key pair mappings

To tackle catastrophic occurrences

- e.g. loss of private keys

To tackle normal exploitation requirements

- e.g. refresh of key pairs for reducing impersonation risks

Problems to solve

Ensure a proper generation of key pairs

Random generation of secret values

- So that they cannot be easily predicted

Increase efficiency without reducing security

- Make security mechanisms more useful
- Increase performance

Goals

Key pair generation

- When and how should they be generated

Handling of private keys

- How do I maintain them private

Distribution of public keys

- How are they correctly distributed worldwide

Lifetime of key pairs

- When will they expire
- Until when should they be used
- How can I check the obsolescence of a key pair

Generation of key pairs: Design principles

Good random generators for producing secrets

Result is indistinguishable from noise

- All values have equal probability
- No patterns resulting from the iteration number or previous values

Example: Bernoulli $\frac{1}{2}$ generator

- Memoryless generator
- $P(b=1) = P(b=0) = \frac{1}{2}$
- Coin toss

Generation of key pairs: Design principles

Facilitate without compromising security

Efficient public keys

- Few 1 bits, typically $2k+1$ values (3, 17, 65537)
- Accelerates operations with public keys
 - Cost is proportional to the number of 1 bits
- No security issues

Generation of key pairs: Design principles

Self-generation of private keys

Maximizes privacy as no other party will be able to use a given private key

- Only the owner has the key
- Even better: The owner doesn't have the key, but may use the key

Principle can be relaxed when not involving signature generation

- Where there are not issues related with non-repudiation

Handling of private keys

Correctness

The private key represents a subject

- e.g., a citizen, a service
- Its compromise must be minimized
- Physically secure backup copies can exist in some cases

The access path to the private key must be controlled

- Access protection with password or PIN
- Correctness of applications that use it

Handling of private keys

Confinement

Protection of the private key inside a (reduced) security domain (ex. cryptographic token)

- The token generates key pairs
- The token exports the public key but never the private key
- The token internally encrypts/decrypts with the private key

Example: SmartCards

- We ask the SmartCard to cipher/decipher something
- The private key never leaves the SmartCard

Distribution of public keys

Distribution to all **senders** of confidential data

- Manual
- Using a shared secret
- Ad-hoc using digital certificates

Distribution to all **receivers** of digital signatures

- Manual
- Ad-hoc using digital certificates

Distribution of public keys

Problem:

How to ensure the correctness of the public key?

Trustworthy dissemination of public keys

- Trust paths / graphs
- If **A trusts K_x^+** , and **B trusts A**, then **B trusts K_x^+**
- Certification hierarchies / graphs
 - With the trust relations expressed between entities
 - Certification is unidirectional!

Public key (digital) certificates

Digital Document issued by a Certification Authority (CA)

Binds a public key to an entity

- Person, server or service

Are public documents

- Do not contain private information, only public one
- Can have additional binding information (URL, Name, email, etc.)

Are cryptographically secure

- Digitally signed by the issuer, cannot be changed

Public key (digital) certificates

Can be used to distribute public keys in a trustworthy way

A certificate receiver can validate it in many ways

- With the CA's public key
- Can also validate the identification
- Validate the validity
- Validate is the key is being properly used

A certificate receiver trusts the behavior of the CA

- Therefore, will trust the documents they sign
- When a CA associates a certificate to A
 - If the receiver trusts the CA
 - Then it will trust that the association of A is correct

Public key (digital) certificates

X.509v3 standard

- Mandatory fields
 - Version
 - Subject
 - Public key
 - Dates (issuing, deadline)
 - Issuer
 - Signature
 - etc.
- Extensions
 - Critical or non-critical

PKCS #6

- Extended-Certificate Syntax Standard

Binary formats

- ASN.1 (Abstract Syntax Notation)
 - DER, CER, BER, etc.
- PKCS #7
 - Cryptographic Message Syntax Standard
- PKCS #12
 - Personal Information Exchange Syntax Standard

Other formats

- PEM (Privacy Enhanced Mail)
- base64 encoding of X.509

Key pair usage

The public certificate binds the key pair to a usage profile

- Private keys are seldom multi-purpose

Typical usage profiles

- Authentication / key distribution
 - Digital signature, Key encipherment, Data encipherment, Key agreement
- Document signing
 - Digital signature, Non-repudiation
- Certificate issuing (exclusively for CAs)
 - Certificate signing, CRL signing
- Timestamping (exclusively for TSAs)

Public key certificates have an extension for this

- Key usage (critical)

Certification Authorities (CA)

Organizations that manage public key certificates

- Companies, not for profit organizations or governmental
- Have the task of validating the relation between key and identity

Define policies and mechanisms for:

- Issuing certificates
- Revoking certificates
- Distributing certificates
- Issuing and distributing the corresponding private keys

Manage certificate revocation lists

- Lists of revoked certificates
- Programmatic interfaces to verify the current state of a certificate

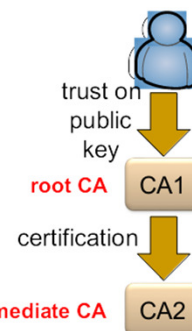
Trusted Certification Authorities

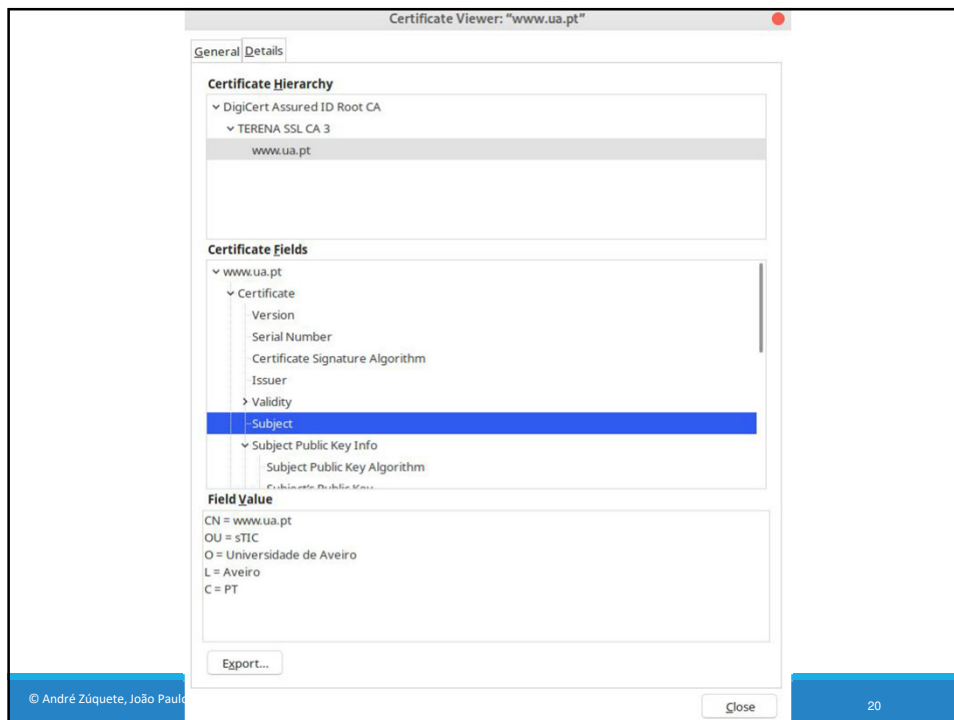
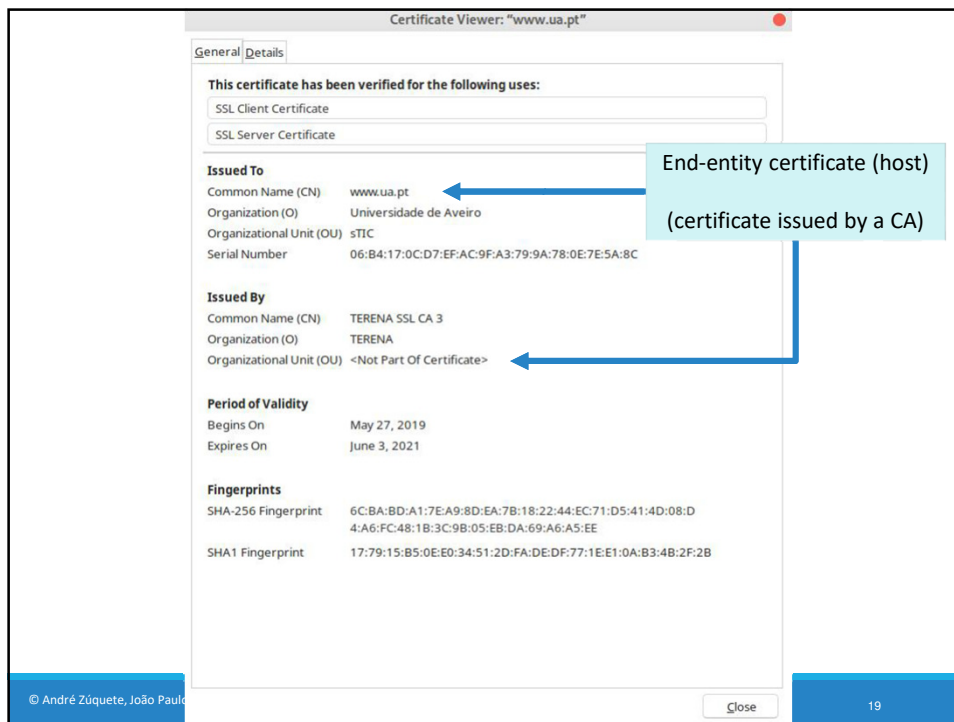
Intermediate CAs: CAs certified by other trusted CAs

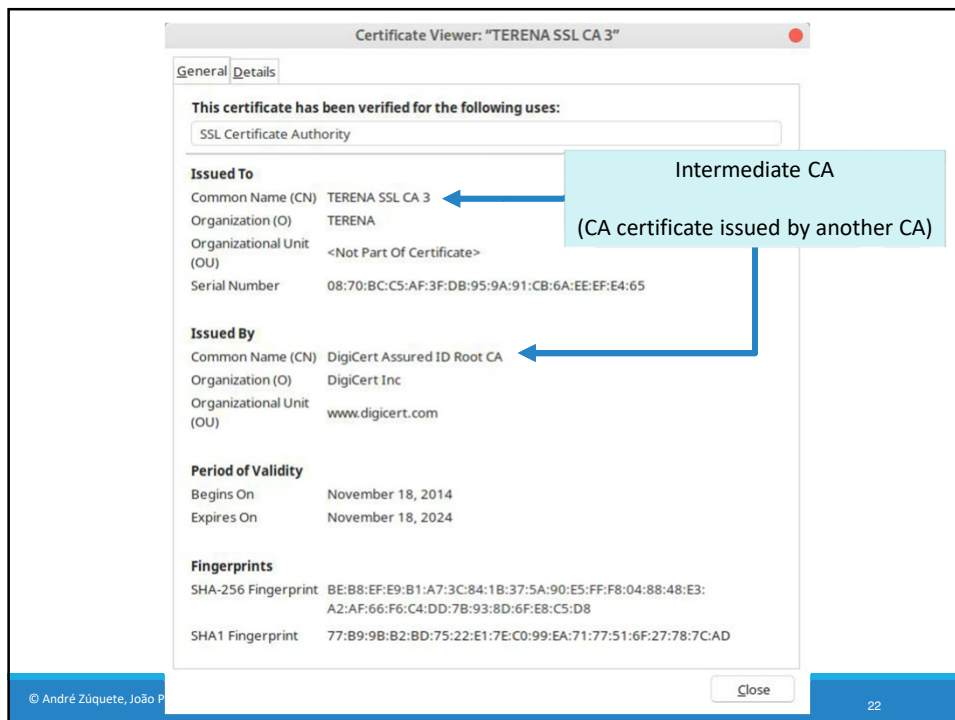
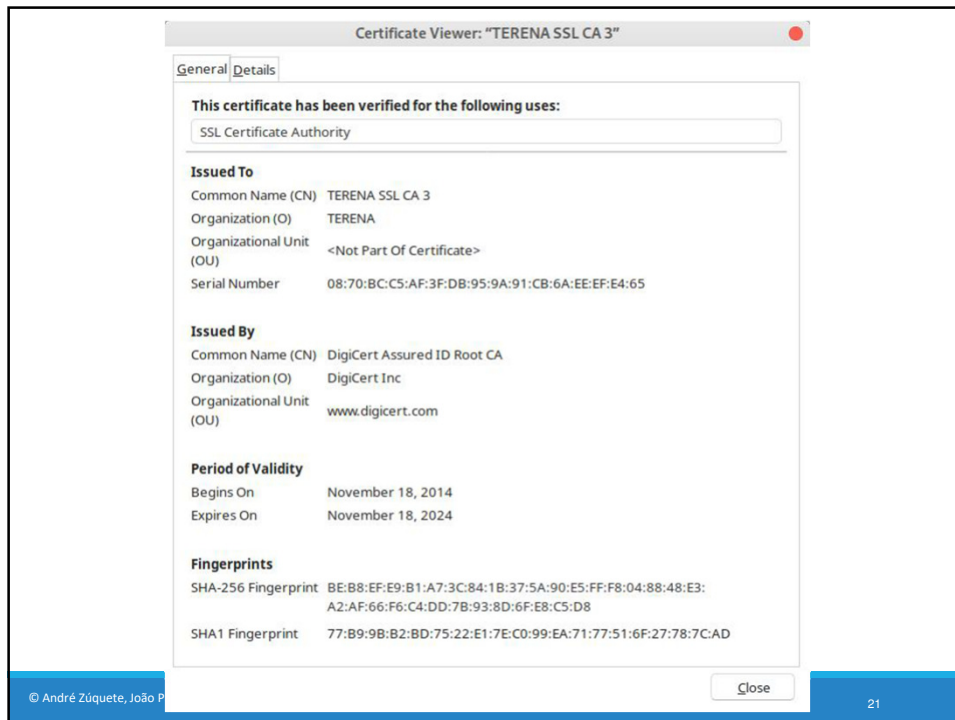
- Using a certificate
- Enable the creation of certification hierarchies

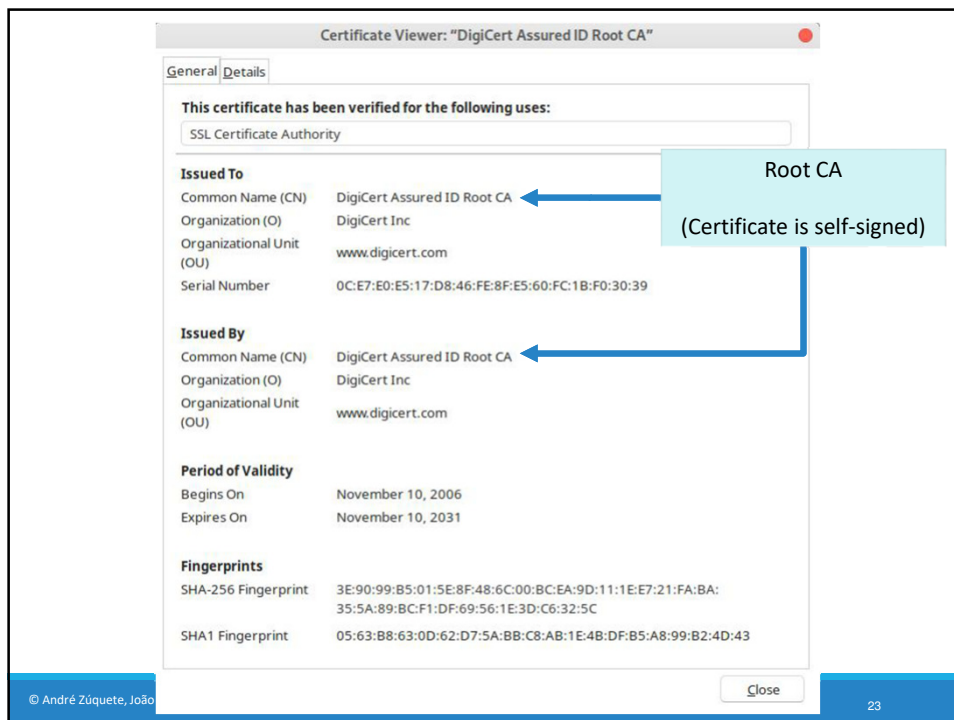
Trusted anchor (or certification root)

- One that has a trusted public key
- Usually implemented by self-certified certificates
 - Issuer = Subject
- Manual distribution
 - e.g., within browsers code (Firefox, Chrome, etc.), OS, distribution...









Refreshing of asymmetric key pairs

Key pairs should have a limited lifetime

- Because private keys can be lost or discovered
- To implement a regular update policy

Problem

- Certificates can be freely copied and distributed
- The universe of holders of certificates is unknown
- Therefore, we cannot contact them to eliminate specific certificates

Solutions

- Certificates with a validity period (not before, not after)
- Certificate revocation lists
 - To revoke certificates before expiring their validity

Certificate revocation lists (CRL)

Base or delta

- Complete / differences

Signed lists of certificates (identifiers) prematurely invalidated

- Must be regularly consulted by certificate holders
- OCSP protocol for single certificate validation
 - RFC 2560
- Can tell the revocation reason

RFC 3280

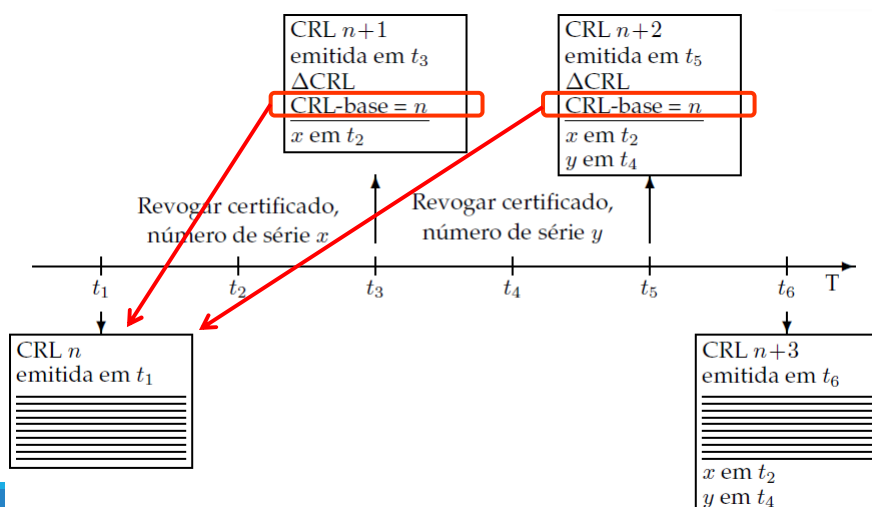
unspecified (0)
keyCompromise (1)
CACompromise (2)
affiliationChanged (3)
superseded (4)
cessationOfOperation (5)
certificateHold (6)

removeFromCRL (8)
privilegeWithdrawn (9)
AACompromise (10)

Publication and distribution of CRLs

- Each CA keeps its CRL and allows public access to it

CRL and Delta CRL



Online Certificate Status Protocol

HTTP-based protocol to assert certificate status

- Request includes the certificate serial number
- Response states if the certificate is revoked
 - Response is signed by the CA and has a validity
- One check per certificate

Requires lower bandwidth to clients

- One check per certificate instead of a bulk download of the CRL

Involves higher bandwidth to CAs

- One check per certificate
- Privacy issues as the CA will know that a certificate is being used

OCSP stapling

- Including a recently signed timestamp in the server response to assert validity
- Reduces verification delay and load on CA
- Avoids privacy issues

Distribution of public key certificates

Transparent (integrated with systems or applications)

- Directory systems
 - Large scale (ex. X.500 through LDAP)
 - Organizational (ex. Windows 2000 Active Directory (AD), Manually (UA IDP))
- On-line: within protocols using certificates for peer authentication
 - eg. secure communication protocols (TLS, IPSec, etc.)
 - eg. digital signatures within MIME mail messages or within documents

Explicit (voluntarily triggered by users)

- User request to a service for getting a required certificate
 - eg. request sent by e-mail
 - eg. access to a personal HTTP page

PKI (Public Key Infrastructure) (1/2)

Infrastructure for enabling a proper use of asymmetric keys and public key certificates

Creation of asymmetric key pairs for each enrolled entity

- Enrolment policies
- Key pair generation policies

Creation and distribution of public key certificates

- Enrolment policies
- Definition of certificate attributes

PKI (Public Key Infrastructure) (2/2)

Definition and use of certification chains (or paths)

- Insertion in a certification hierarchy
- Certification of other CAs

Update, publication and consultation of CRLs

- Policies for revoking certificates
- CRL distribution services
- OCSP services

Use of data structures and protocols enabling inter-operation among components / services / people

PKI Example: Citizen Card

Enrollment

- In loco, personal enrolment

Multiple key pairs per person

- One for authentication
- One for signing data
- Both generated inside smartcard, not exportable
- Both require a PIN to be used in each operation

Certificate usage (authorized)

- Authentication
 - SSL Client Certificate, Email (Netscape cert. type)
 - Signing, Key Agreement (key usage)
- Signature
 - Email (Netscape cert. type)
 - Non-repudiation (key usage)

Certification path

- Uses a well-known, widely distributed root certificate
 - GTE Cyber Trust Global Root
- PT root CA below GTE
- CC root CA below PT root CA
- CC Authentication CA and CC signature CA below CC root CA

CRLs

- Signature certificate revoked by default
 - Revocation is removed if the CC owner explicitly requires the usage of CC digital signatures
- All certificates are revoked upon a owner request
 - Requires a revocation PIN
- CRL distribution points explicitly mentioned in each certificate

Certificate Pinning

If attacker has access to trusted Root, it can impersonate every entity

- Manipulate a trusted CA into issuing certificate (unlikely)
- Inject custom CA certificates in the victim's database (likely)

Certificate Pinning: add the fingerprint of the PubK to the **source code**

- Fingerprint is a hash (e.g. SHA256)

Validation process:

- Certificate must be valid according to local rules
- Certificate must have a public key with the given fingerprint

Certification Transparency (RFC 6962)

Problems

- CAs can be compromised (e.g., DigiNotar)
 - By attackers
 - By governments, etc.
- Compromise is difficult to detect
 - Result in the change of assumptions associated to the behavior of the CA
 - Owner will seldom know

Definition: a global system records all public certificates created

- Ensure that only a single certificate has the correct roots
- Stores the entire certification chain of each certificate
- Presents this information for auditing
 - Organizations or ad-hoc by the end users

Authentication Mechanisms and Protocols

Authentication (Authn)

Proof that an entity has an attribute it claims to have

- Hi, I'm Joe
- Prove it!
- Here is my **proof**,
calculated with **Joe's credentials** that I've agreed with you
- Proof accepted/not accepted

- Hi, I'm over 18
- Prove it!
- Here is a **claim** issued by a competent authority,
which I can also **prove** that I'm the owner
- Proof and claim accepted/not accepted

Authn: Proof Types

Something we know

- A secret memorized (or written down...) by Joe

Something we have

- An object/token solely held by Joe

Something we are

- Joe's Biometry

Multi-factor authentication

- Simultaneous use of different proof types
- 2FA = Two Factor Authentication

Risk-based MFA

- Variable MFA
- Higher attack risk, more factors or less risky factors
- Lower attack risk, less or easier factors

Authn : Goals

Authenticate interactors

- People, services, servers, hosts, networks, etc.

Enable the enforcement of authorization policies and mechanisms

- Authorization \neq authentication
- Authorization \Rightarrow authentication

Facilitate the exploitation of other security-related protocols

- e.g. key distribution for secure communication

Authn : Requirements

Trustworthiness

- How good is it in proving the identity of an entity?
- How difficult is it to be deceived?
- Level of Assurance (LoA)

Secrecy

- No disclosure of secret credentials used by legit entities

NIST 800-63				
LoA	DESCRIPTION	TECHNICAL REQUIREMENTS		
		IDENTITY PROOFING REQUIREMENTS	TOKEN (SECRET) REQUIREMENTS	AUTHENTICATION PROTECTION MECHANISMS REQUIREMENTS
1	Little or no confidence exists in the asserted identity; usually self-asserted; essentially a persistent identifier	Requires no identity proofing	Allows any type of token including a simple PIN	Little effort to protect session from off-line attacks or eavesdropper is required.
2	Confidence exists that the asserted identity is accurate; used frequently for self service applications	Requires some identity proofing	Allows single-factor authentication. Passwords are the norm at this level.	On-line guessing, replay and eavesdropping attacks are prevented using FIPS 140-2 approved cryptographic techniques.
3	High confidence in the asserted identity's accuracy; used to access restricted data	Requires stringent identity proofing	Multi-factor authentication , typically a password or biometric factor used in combination with a 1) software token, 2) hardware token, or 3) one-time password device token	On-line guessing, replay, eavesdropper, impersonation and man-in-the-middle attack are prevented. Cryptography must be validated at FIPS 140-2 Level 1 overall with Level 2 validation for physical security.
4	Very high confidence in the asserted identity's accuracy; used to access highly restricted data.	Requires in-person registration	Multi-factor authentication with a hardware crypto token.	On-line guessing, replay, eavesdropper, impersonation, man-in-the-middle, and session hijacking attacks are prevented. Cryptography in the hardware token must be validated at FIPS 140-2 level 2 overall, with level 3 validation for physical security.

Authn : Requirements

Robustness

- Prevent attacks to the protocol data exchanges
- Prevent on-line DoS attack scenarios
- Prevent off-line dictionary attacks

Simplicity

- It should be as simple as possible to prevent entities from choosing dangerous shortcuts

Deal with vulnerabilities introduced by people

- They have a natural tendency to facilitate or to take shortcuts
- Deal with phishing!

Authn: Entities and deployment model

Entities

People

Hosts

Networks

Services / servers

Deployment model

Along the time

- Only when interaction starts
- Continuously along the interaction

Directionality

- Unidirectional
- Bidirectional (Mutual)

Authn interactions: Basic approaches

Direct approach

1. Provide credentials
2. Wait for verdict

- Advantage: **no computations** by the presenter
- Disadvantage: credentials can be **exposed** to malicious validators

Challenge-response approach

1. Get challenge
2. Provide a response computed from the challenge and the credentials
3. Wait for verdict

- Advantage: credentials are **not exposed** to malicious validators
- Disadvantage: requires **computations** by the presenter

Authn of subjects: Direct approach w/ known password

A password is checked against a value previously stored

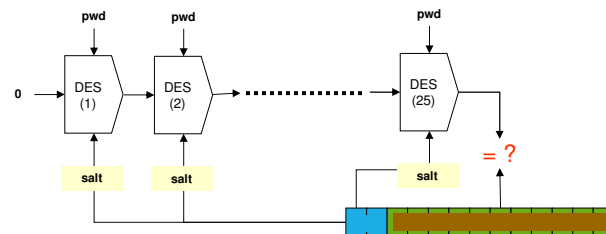
- For a claimed identity (username)

Personal stored value:

- Transformed by a unidirectional function
- Windows: digest function
- UNIX: DES hash + salt
- Linux: MD5 + salt
 - hash is configurable

Optimal: PBKDF2, Script with high complexity

Authentication of subjects: Direct approach w/ known password



$DES_{hash} = DES_{pwd}^{25}(0)$
 $DES_k^n(x) = DES_k(DES_k^{n-1}(x))$
 Permutation of 12 subkeys' bit pairs with salt (12 bits)

Authn of subjects: Direct approach w/ known password

Advantage

- Simplicity!

Problems

- Usage of weak passwords
 - Enable dictionary attacks
- Transmission of passwords along insecure communication channels
 - Eavesdroppers can easily learn the password
 - e.g. Unix remote services, PAP

Top Ten 2017 from Splashdata

1. 123456
2. Password
3. 12345678
4. qwerty
5. 12345
6. 123456789
7. letmein
8. 1234567
9. football
10. iloveyou

Authn of people: Direct approach with biometrics

People get authenticated using body measures

- Biometric samples
- Fingerprint, iris, face geometry, voice timber, manual writing, vein matching, etc.

Measures are compared with personal records

- Biometric references (or template)
- Registered in the system with a previous enrolment procedure

Identification vs authentication

- Identification: 1-to-many check for a match
- Authentication: 1-to-1 check for a match



Authn of people: Direct approach with biometrics

Advantages

- People do not need to use memory, or carry something
 - Just be their self
- People cannot choose weak passwords
 - In fact, they don't choose anything
- Authentication credentials cannot be transferred to others
 - One cannot delegate its own authentication

Authentication of people: Direct approach with biometrics

Problems

- Biometric methods are still incipient
 - In many cases it can be fooled with ease (Face Recognition, Fingerprint)
- People cannot change credentials
 - If the credentials or templates are stolen
- Credentials cannot be transferred between individuals
 - If it is required in extraordinary scenarios
- Can pose risks to individuals
 - Physical integrity can be compromised by an attacker in order to acquire biometric data
- It is not easy to be implemented in remote systems
 - It is mandatory to have secure and trusted biometric acquisition devices
- Biometrics can reveal other personal secrets
 - Diseases

Authn of subjects:
Direct approach with one-time passwords

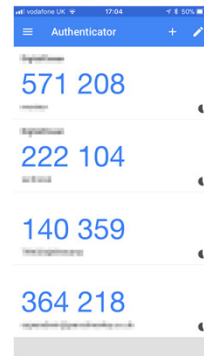
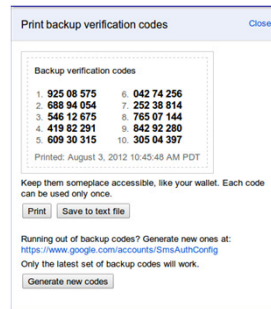
One-Time Passwords = Secrets that can be used only once

- Pre-distributed directly, or the result of a generator function

Example: Bank codes, Google Backup Codes



https://www.montepio.pt/SitePublico/pt_PT/particulares/montepio24/cartao-matriz.page?altcode=10006P



Authn of subjects:
Direct approach with one-time passwords

Advantages

- Can be eavesdropped, allowing its use in channels without encryption
- Can be chosen by the authenticator, which may enforce a given complexity
- Can depend on a shared password

Problems

- Interacting entities need to know which password to use on each occasion
 - Implies some form of synchronization (e.g., index, coordinates)
- Individuals may require additional resources to store/generate the passwords
 - Sheet of paper, application, additional device, etc.

Yubikey

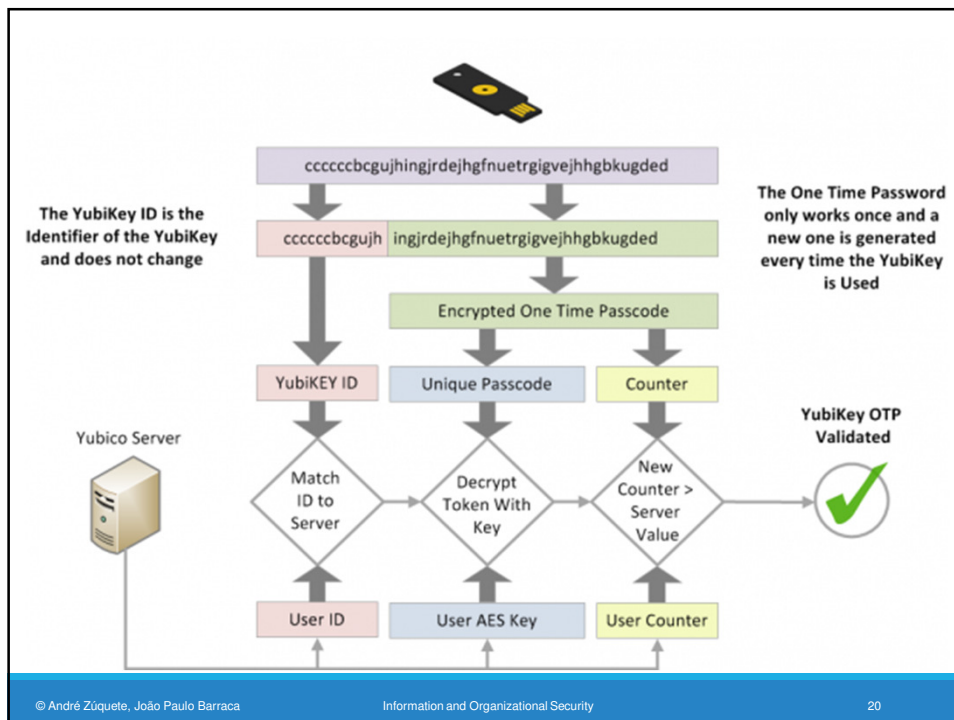
Personal Authentication Device

- USB, Bluetooth and/or NFC



Activation generates a 44 characters key

- Emulates a USB keyboard (besides an own API)
- Supports HOTP (events) or TOPT (Temporal)
- If a challenge is provided, user must touch the button to obtain a result
- Several algorithms, including AES 256



Challenge-response Approach

The authenticator provides a challenge

- A nonce (value not once used)
- Usually random
- Can be a counter

The authenticated entity transforms the challenge

- The transformation method is shared with the authenticator

The result is sent to the authenticator

The authenticator verifies the result

- Calculates a result using the same method and challenge
- Or produces a value from the result and evaluates if it is equal to the challenge, or to some related value

Challenge-Response Approach

Advantages

- Authentication credentials are not exposed
- An eavesdropper will see the challenge and the result
 - but has no knowledge about the transformation

Problems

- Authenticated entities must have the capability of calculating results to challenges
 - Hardware token ou software application
- The authenticator may need to keep shared secrets (in clear text)
 - Secrets can be stolen
 - Individuals may reuse secrets in other systems, enabling lateral attacks
- May be possible to calculate all results to a single (or all) challenge(s)
 - Can reveal the secret used
- May be vulnerable to dictionary attacks
- Authenticator should NEVER issue the same challenge to the same user

Authn of Subjects: Challenge-Response with Smartcards

Authentication Credentials

- Having the smartcard
 - e.g., the Citizen Card
- The private key stored inside the smartcard
- The PIN code to access the key

The authenticator knows

- The user public key

Robust against:

- Dictionary attacks
- Offline attacks to the database
- Insecure channels

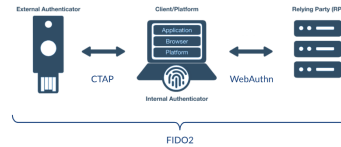


Authn of Subjects: Challenge-Response with Smartcards

Challenge-Response Protocol

- The authenticator generates a challenge
- Smartcard owner ciphers the challenge with their private key
 - Stored in the smartcard, protected by the PIN code
 - In alternative, can sign the challenge
- The authenticator deciphers the result with the public key
 - If the decrypted result matches the challenge, the authentication is successful
 - In alternative, it can verify the signature (which is the same process)

Authn of Subjects: Challenge-Response with other tokens



FIDO2 tokens (FIDO Alliance)

- For both mobile and desktop environments
- Web Authentication (WebAuthn) specification
- Client-to-Authenticator Protocol (CTAP)
- Security
 - Credentials never leave the user's device and are never stored on a server
 - No risks of phishing, no password theft (still, tokens can be stolen)
 - No replay attacks
 - Token certification levels
- Privacy
 - Credentials are unique per website
 - Tracking is not possible (different web sites, different public keys for the same token)
 - Biometric data, when used, never leaves the user's device

<https://www.inovex.de/de/blog/fido2-webauthn-in-practice/>

FIDO2 certification

FIDO Authenticator Certification Examples

L3+		USB U2F Token built on a CC-certified Secure Element Certification: L3+
L3		USB U2F Token built on a basic simple CPU, OS, is certified. Good physical anti-tampering enclosure
L2		UAF implemented as a TA in an uncertified TEE
L1+		UAF in downloadable app using white box crypto and other techniques Certification: L1+
L1		Downloaded app making use of Touch ID on iOS Certification: L1
		FIDO2 making use of the Android keystore. Keystore is not certified Certification: L1
		FIDO2 built into a downloadable web browser app Certification: L1

Authn of Subjects: Challenge-Response with Shared Secret

Authentication Credentials

- Password selected by the individual

The authenticator knows:

- Bad approach: the shared password
- Better approach: A transformation of the shared password
 - The transformation should be unidirectional

Authentication of Subjects: Challenge-Response with Shared Secret

Basic Challenge-Response Protocol

- The authenticator generates a challenge
- The individual calculates a transformation of the challenge and the password
 - $\text{result} = \text{hash}(\text{challenge} || \text{password})$
 - or... $\text{result} = \text{encrypt}(\text{challenge}, \text{password})$
- The authenticator reverts the process and checks if the values match
 - $\text{result} == \text{hash}(\text{challenge} || \text{password})$
 - or $\text{challenge} == \text{decrypt}(\text{result}, \text{password})$
- Examples with shared passwords: CHAP, MS-CHAP v1/v2, S/Key
- Examples with shared keys: SIM & USIM (celular communications)

PAP and CHAP (RFC 1334, 1992, RFC 1994, 1996)

Protocols user for PPP (Point-to-Point Protocol)

- Unidirectional authentication
 - The authenticator authenticates users, but users do not authenticate the authenticator

PAP (PPP Authentication Protocol)

- Simple presentation of a UID/password pair
- Insecure transmission (in clear text)

CHAP (CHallenge-response Authentication Protocol)

Aut → U : authID, challenge

U → Aut: authID, MD5(authID, secret, challenge), identity

Aut → U : authID, OK/not OK

- The authenticator can request further authentication at any time

Authentication of subjects: Challenge-Response with Shared Key

Uses a cryptographic key instead of a password

- Robust against dictionary attacks
- Requires a device to store the shared key

GSM Subscriber authentication

Uses a secret shared between the HLR and the subscriber phone

- Uses 128-bit shared key (not an asymmetric key pair)
- Key is stored in the SIM card
- SIM card is unlocked by a user PIN
- SIM card answers challenges using the shared key

Uses (initially unknown algorithms):

- A3 for authentication
- A8 to generate the session key
- A5 is a stream cipher for communication

A3 and A8 executed by the SIM, A5 executed by the baseband

- A3 and A8 can be chosen by the operator

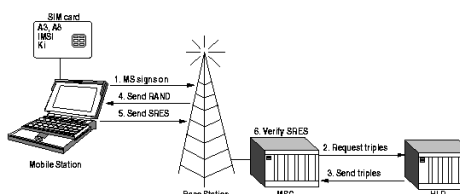
GSM Subscriber authentication

MSC requests triples from HLR/AUC

- **RAND, SRES, Kc**
- It can ask one or several

HLR generates RAND and the triples using the subscriber Ki

- **RAND**, random value (128 bits)
- **SRES = A3 (Ki, RAND)** (32 bits)
- **Kc = A8 (Ki, RAND)** (64 bits)



Frequently uses COMP128 for the A3/A8 algorithms

- Recommended by the GSM consortium
- **[SRES, Kc] = COMP128 (Ki, RAND)**

Authentication of Systems

By name (DNS) or MAC/IP address

- Extremely weak, without cryptographic proof
- Still... it is used by some services
- e.g., NFS, TCP wrappers

With cryptographic keys

- Secret keys, shared between entities that communicate frequently
- Asymmetric key pairs, one per host
 - Public keys pre-shared with entities that communicate frequently
 - Public keys certified by a third party (a CA)

Authentication of Services

Authentication of the host

- All services co-located in the same host are automatically and indirectly authenticated

Credentials exclusive to each service

Authentication:

- Secret keys shared with clients
 - When they require authentication of the clients
- Asymmetric key pairs by host/service
 - Certified by others or not

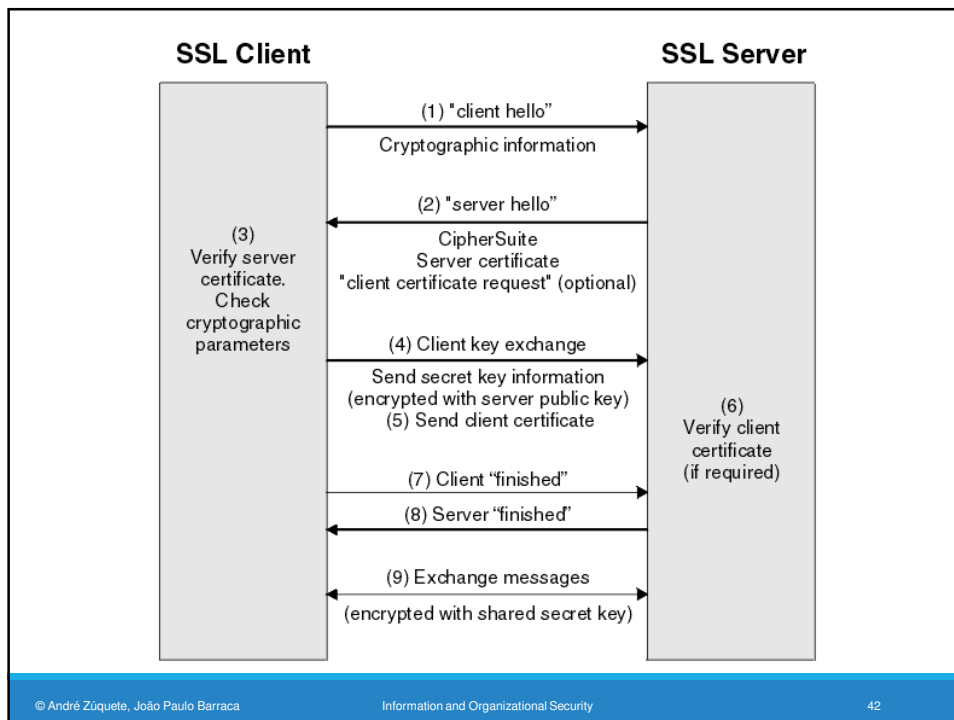
TLS (Transport Layer Security, RFC 2246)

Secure Communication Protocol over TCP/IP

- Evolved from the SSL V3 (Secure Sockets Layer) standard
- Manages secure sessions over TCP/IP, individual to each application
 - Initially designed for HTTP traffic
 - Currently used for many other types of traffic

Security mechanisms

- Confidentiality and integrity of the communication between entities
 - Key distribution, negotiation of ciphers, digests and other mechanisms
- Authentication of the intervenient entities
 - Servers, services, etc...
 - Clients (not so common)
 - Both executed with asymmetric keys and X.509 certificates



TLS Ciphersuites

If a server supports a single algorithm, it cannot be expected for all clients to also support it

- More powerful/limited, older/newer

The Ciphersuite concept allows the negotiation of mechanisms between client and server

- Both send their supported ciphersuites, and select one they both share
- The server chooses

Exemplo: ECDHE-RSA-AES128-GCM-SHA256

Format:

- Key negotiation algorithm: ECDHE (Elliptic Curve Ephemeral Diffie-Hellman)
- Authentication algorithm: RSA
- Cipher algorithm and cipher mode: AES-128 GCM
- Integrity control algorithm: SHA256

SSH (Secure SHell)

Manages secure console sessions over TCP/IP

- Initially designed to replace the Telnet application/protocol
- Currently used in many other applications
 - Execution of remote commands in a secure manner (rsh/rexec)
 - Secure copy of contents from/to remote hosts (rcp)
 - Secure FTP (sftp)
 - Secure (Generic) communication tunnels (carry standard IP packets)

Security Mechanisms

- Confidentiality and integrity of the communications
 - Key distribution
- Authentication of the intervening entities
 - Server / Hosts
 - Client users
 - Both achieved through several, and differentiated mechanisms

SSH: Authentication Mechanisms

Server: a pair of asymmetric keys

- Keys are distributed during the interaction
 - Not certified!
- Clients store the public keys from previous interactions
 - Key should be stored in some trusted environment
 - If the key changes the client is warned
 - e.g., server is reinstalled, key is regenerated, an attacker is hijacking the connection
 - Client can refuse to continue with the authentication process

Clients: authentication is configurable

- Default: username and password
- Other: username + private key
 - The public key MUST be pre-installed in the server
- Other: integration with PAM for alternative authentication mechanisms

Centralized network authentication

Used for restricting network access to known clients

- In cabled networks
- In wireless networks
- In VPNs (Virtual Private Networks)

Usually implemented by a central service

- AAA server
 - Authentication, Authorization and Accounting
 - e.g. RADIUS and DIAMETER
- This server defines which network services the user can make use of

Authentication by an IdP

Unique, centralized authentication for a set of federated services

- The identity of a client, upon authentication, is given to all federated services
- The identity attributes given to each service may vary
- The authenticator is called **Identity Provider (IdP)**
- The federated service is called a **Relying Party (RP)**
- In some cases, the provided identity attributes are shown to the client

Examples

- Authentication at UA
 - Performed by a central, institutional IdP (idp.ua.pt)
 - The identity attributes are securely conveyed to the service accessed by the user
- Autenticação.gov (www.autenticacao.gov.pt)
 - Performed by a central, national IdP
 - The identity attributes are shown to the user
- Other:
 - Services used worldwide: Google, Facebook, etc.

Centralized authentication

Advantages:

- Can reuse same credentials over multiple systems/services
- Single secure repository for credentials
 - More difficult to steal credentials when used in many services
- Can implement restrictions to services/systems

Disadvantages:

- Requires additional servers
- Single point of failure: without authentication systems, no one will be authenticated
 - Important to also deploy local credentials for admins
- Introduces delays in the authentication process

Single Sign-On

A facility usually associated with IdP

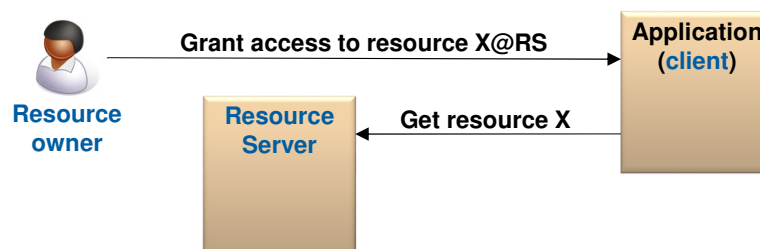
- Both not mandatory nor always appropriate

SSO exists for simplifying users' life

- They login just one for accessing several federated services during a given time period

OAuth 2.0: delegation (RFC 6749)

A framework to allow users to delegate access to their resources on their behalf



OAuth 2.0 roles

Resource owner

- An entity capable of granting access to a **protected resource**
- **End-user**: a resource owner that is a person

Client

- An **application** making requests for protected resources on behalf of the resource owner and with its authorization

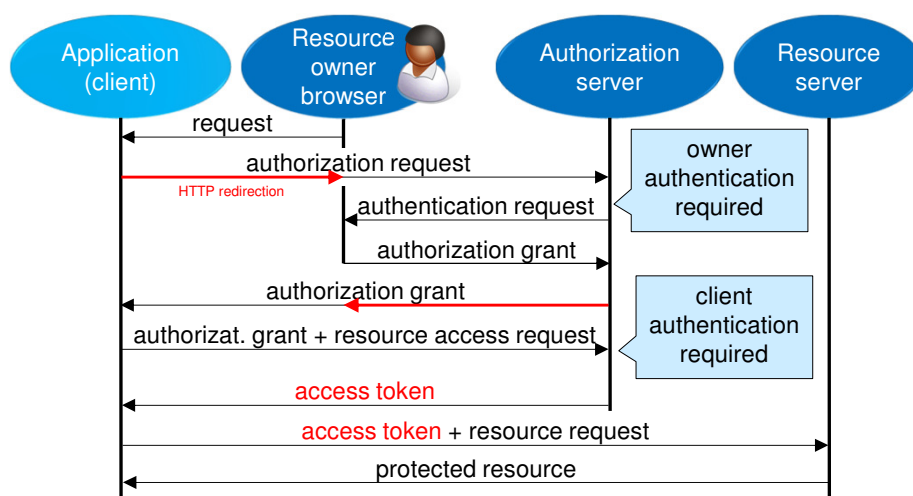
Resource Server

- The server hosting protected resources
- Responds to protected resource requests using **access tokens**

Authorization Server

- The server issuing access tokens to clients after successfully **authenticating resource owners** and obtaining their **authorization** for the clients to access one of their (users) resources

Protocol flow



OpenID Connect (OIDC)

An identification layer on top of OAuth 2.0

- OAuth 2.0 provides the fundamental centralized authentication
- The protected resources are identity attributes
 - Packed in **scopes**
 - The attributes are called (identity) **claims**

Reliable storage

Problems

Storage devices develop faults

- It should be minimized the failures in storage devices and loss of data
- Failure is certain and cannot be ignored

Access to mechanical disks is slow (hard disks)

- Access Time = Translation time + Rotation Time
- More information → higher impact of storage media

Problems

Solid State Devices (SSDs) have a limited number of write operations

- 2000-3000 writes per sector for MLC (2 bits per cell)

Specific events may result in total data loss

- Fire, robbery, “energy peaks”, floods, user mistakes, attacks

May be required to distribute data in an intelligent manner

- To maximize performance
- To reduce costs

Solutions

Data backups

- Local
- Remote

Redundant Storage

- RAID
- Other: ZFS

Better storage devices, environments with higher control

- SLED (Single Large Expensive Disks)
- Enterprise Grade devices
- Temperature and Humidity Control

Infrastructures dedicated for storage

- Single policy control point

Backups

Periodic copy of data

- Snapshot of the storage state in a specific moment
- Copies will allow to set files to a previous version
- May be encrypted

Full: Complete snapshot of the data volume

- Fast recovery
- Requires a large amount of space

Differential: Differences since the last full backup

- Slower recovery, but also lower storage requirements
- Daily differential backups will grow as changes increase

Incremental: Differences since the last backup

- Even slower recovery
- Requires reconstruction of all intermediate backups since the last full
- Higher storage space efficiency

Backups

A backup is not an additional disk with data

- External or remote

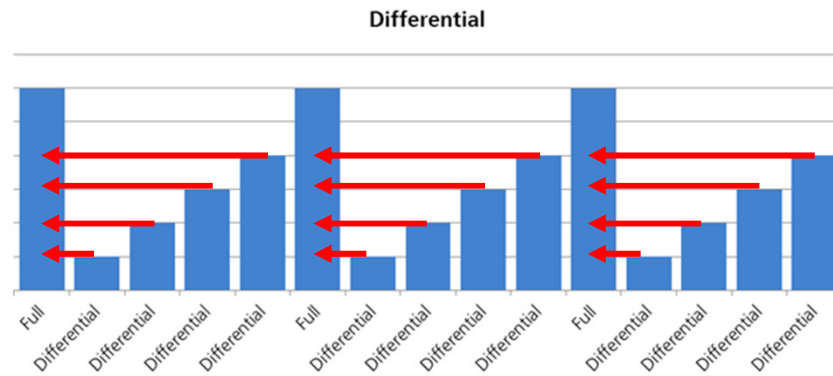
It considers policies, mechanisms and processes to make, maintain and recover copies of the same data

- Should resist specific situations
- Should be used only in emergency situations
- Important to consider both the copy, storage and recovery!

Legal framework implies a special care

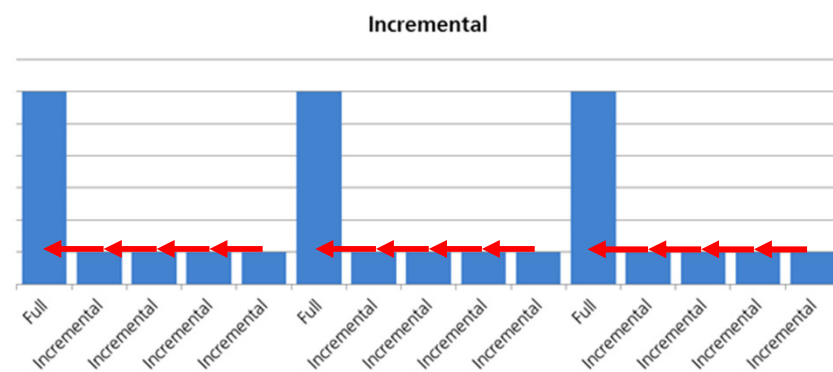
- When dealing with personal data
- Frequently impose a retention policy
 - Backups should expire after some time

Backups types: Differential



<http://www.teamlead.co.uk/>

Backups types: Incremental



<http://www.teamlead.co.uk/>

Backups: Compression

Uses lossless compression algorithms and solutions

- Ex: ZIP

Copy only some parts of the information

- Only modified files

Deduplication

- Only store unique files/blocks
- Usually using full copy with offline deduplication
 - Of disk blocks using specific image formats
 - Of files using hard links

Backups: Levels

Applications

- Extract data from applications (e.g. mysqldump)
- Represent a consistent view of the application
 - May be required to block the application state (e.g., database changes)
- May be repeated for each individual application

Files

- Copy of individual files
- May backup any application in a filesystem
- State may be inconsistent
 - e.g., open files without data written, or applications change many files at once

Backups: Levels

Filesystem

- Internal features provided by each individual filesystem
- Creation of periodic snapshots with records of all changes or current state
- May allow the recovery of individual files, or the entire filesystem

Device Blocks

- Copy of all blocks of a storage medium
- Independent of the filesystem or operation system in use
- May be implemented by the storage infrastructure
 - Transparent and without any impact to applications

Backups: Location of data

In the same volume or in the same server

- Allow users to rapidly recover information
- Protects against changes/deletions made by users
- May not protect against hardware malfunction
 - e.g., macOS Timemachine

In a system location in the same infrastructure

- Also, with fast access time
- Protects against isolated storage failures
- Doesn't protect data against events with broader reach
 - Floods, fire, robbery
- Examples: Most enterprise storage solutions, backuppc, TimeCapsule, Borg, Kopia

Backups: Location of data

Remote (off-site)

- Implemented to a system outside the local datacenter
 - Dedicated service or through the internet
 - e.g., Amazon S3, or to servers in a dedicated datacenter
 - Encryption if recommended (or mandatory) in the case of external services!
- Implemented with specialized secure transport
 - Armored car transporting backups to a secure place
- Allow recovery even if far reaching events occur
 - Terrorism, Earthquake
- Recovery will be slower
 - Limited by the speed of a network link or the physical transport

Selecting Storage Devices

Different device grades: Enterprise vs Desktop

- Different construction quality and recovery features
- Different MTBF: Mean Time Between Failures
 - Enterprise HDD: 1.2M hours, at 45°C, working 24/7, 100% use rate (1)
 - Desktop HDD: 700K hours, at 25°C, working 8/5, 10-20% use rate(1)

Adjusted to each use case

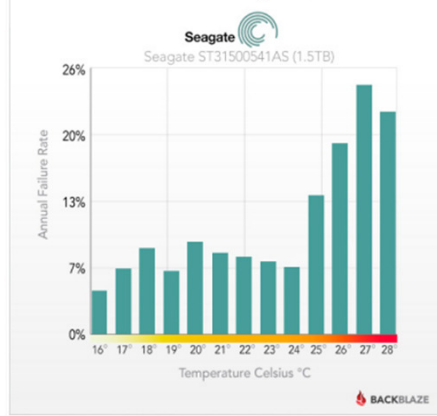
- Write intensive vs Read Intensive
- NAS vs Video vs Desktop vs Cold Storage vs Data Center
 - Differences in power consumption, reliability and performance

Adjusted to a specific performance level

- Tier 0: Highest performance, low capacity (PCIe NVME SLC SSD)
- Tier 1: Some performance, high capacity and availability (M2 SATA SSD)
- Tier 3: Low performance, high capacity, low price (SATA HDD)

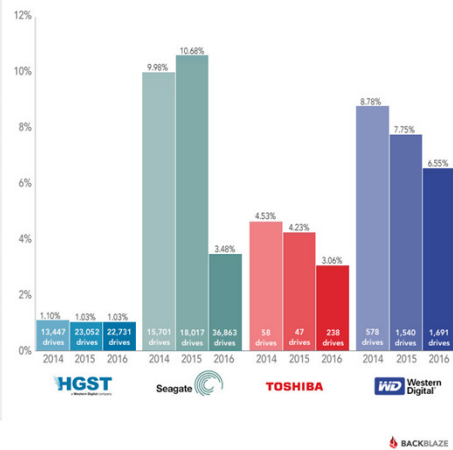
Controlled Environment and Equipment

Failure Rate of a Seagate Drive



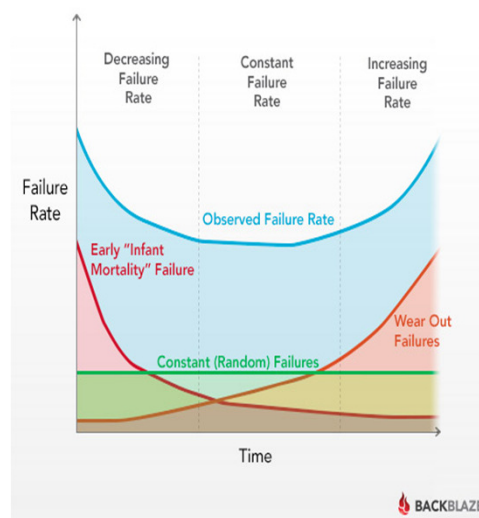
Hard Drive Failure Rates by Manufacturer

All drive sizes for a given Manufacturer are combined



<https://www.backblaze.com/b2/hard-drive-test-data.html>

Controlled Environment and Equipment



RAID: Redundant Array of Inexpensive Drives

Improves the survivability of information

- Data is only lost after several devices are lost
- The number of lost devices is configurable

Low cost and efficient solution

- Can use cheap, lower quality hardware
- Can improve read and write performance

RAID doesn't replace backups

- Only tolerates the failure of a limited number of devices
- Cannot cope with user mistakes (file modification/deletion)

RAID can even increase the failure probability

- As it can be tweaked towards performance

RAID 0 (Striping)

Objectives

- Speedup data access

Approach

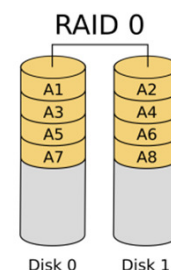
- Access disks in parallel
- Striping
 - Data is split in small chunks (stripes)
 - Stripes are stored among all disks in a distributed manner

Advantages

- May speedup performance as a factor of the number of disks

Disadvantages

- Increases the probability of losing data
 - If P_f is the probability of failure of a single disk, an N -disk RAID 0 volume will have a $1-(1-P_f)^N$ failure probability
- Increases the number of devices
 - At least it will double the number



RAID 1 (Mirroring)

Objectives

- Tolerate disk failures

Approach

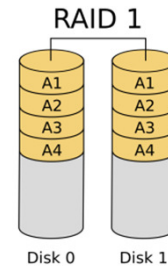
- Data duplication (mirroring)
- Synchronized writing
- Distributed read from any disk with or without comparison from another disk

Advantages

- Decreases the probability of data loss
- If P_f is the probability of failure of a single disk, the probability of failure with N disks is P_f^N

Disadvantages

- Storage inefficiency
- Will lose at least 50% of the total capacity
- For 3 disks it will lose 66%... Loss is $(N-1)/N$
- Increase the number of devices
- At least to the double



RAID 0+1 and 1+0 (Nested)

Objectives

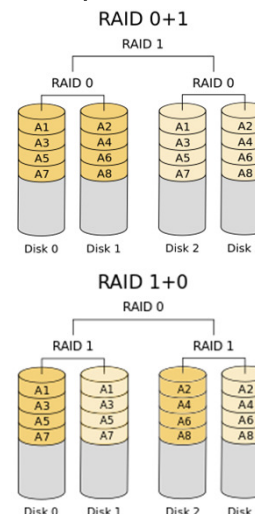
- Benefits of RAID 0 (performance)
- Benefits of RAID 1 (resilience)

Approach

- 0+1: A RAID 1 volume using RAID 0 volumes
- Mirroring of striped volumes
- 1+0: RAID 0 over RAID 1 volumes
- Striping over mirrored volumes

Disadvantages

- Storage capacity waste
- At least 50%
- Increase the number of devices



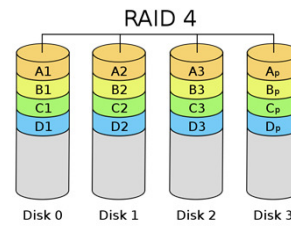
RAID 4

Objectives

- Have some resilience as RAID 1
- With a performance close to RAID 0

Approach

- Store data in N-1 disks
- Store parity data in an additional disk
- Total waste is dependent on the capacity and number of disks
- Data from any N-1 disk can be used to recreate another one



Disadvantages

- Requires at least 3 disks
- Updating parity data is complex and will require specific hardware
- Imposes the need to read before any write
 - Read data from existing block (e.g., C1) and from the corresponding parity disk (Cp)
 - Compare old data block with new, and change the parity block (Cp')
 - Write the new data block (C1') and the new parity block (Cp')
- Writes must be serialized due to the existence of a parity disk
- Recovery is way more complex than with RAID 1

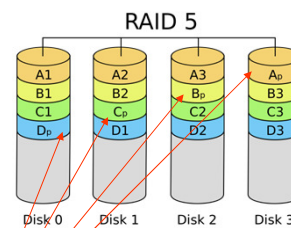
RAID 5

Objectives

- Similar to RAID 4
- But with higher write efficiency

Approach

- Distribute the parity blocks among all disks
- Waste is similar to RAID 4
- Write concurrency is improved



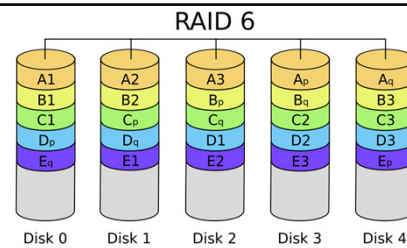
Disadvantages

- More complex to be implemented

RAID 6

Objectives

- Improve the reliability of RAID 5



Approach

- Use 2 parity blocks, distributed among all disks
- Capacity waste will be higher than in RAID 5 (equal to 2 disks)
- Concurrency is slightly worse than with RAID 5

Advantages

- Allows the failure of two disks without data loss

Disadvantages

- Even more complex than RAID 5

NAS and SAN

NAS: Network Attached Storage

- Storage system available in the network
- Frequently created with RAID disks
- Cost: Hundreds to Thousands of Euro

SAN: Storage Area Network

- Set of systems available in a network
- Implemented distributed storage with redundancy
- Cost: Hundreds of Thousands to Millions of Euro

Advantages

- Allow centralizing the storage policies
- Provide a normalized interface, independent of the real storage
- May be used to distributed backups

Confidential data storage

Problems

The protections provided by a traditional filesystem are limited

Physical Protections

- File system is limited to a physical device

Logical Protections

- Access control to files, controlled by the operating system
- Using ACLs and other confinement mechanisms

Problems

There is a relevant number of situations where standard protections are irrelevant

When there is direct and physical access to devices

- Access to host devices (laptops, smartphones, servers)
- Access to external storage devices
 - Tapes, CDs, DVDs, SSDs, NAS

Access through the system with the correct rights

- Non-ethical access by system administrators
- With impersonation attacks

Problems

**There is a prevalence of distributed storage
It imposes trusting multiple administrators, sometimes unknown**

Authentication is made remotely

- Sometimes it is not clear what is the security level of said methods
- Storage Provider may have unknown integrations
- Interaction models are complex, through external networks
- Multiple entities involved

Information is transmitted through communication channels

- May violate confidentiality, integrity and create privacy issues

Solution: Encrypt data

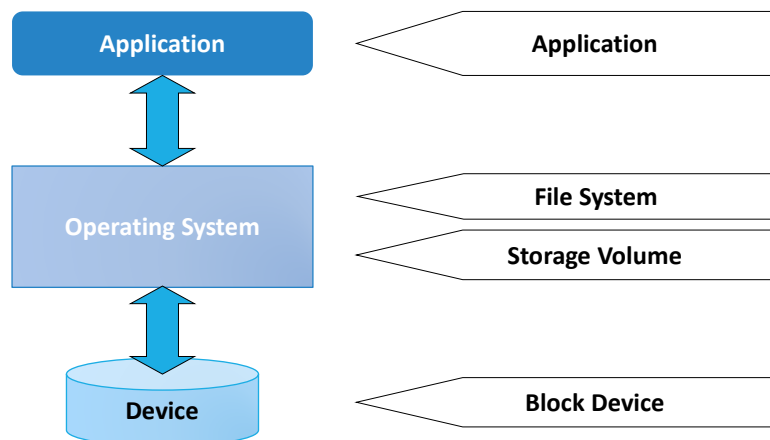
Encryption/Decryption of file contents

- Enable secure transfer over insecure networks
- Enable secure storage in insecure locations
 - Managed by external entities, or in shared storages

Problems of encryption

- Access to information
 - Users may lose the keys
 - Key loss = data loss
 - Key storage may reduce overall security
- File sharing
 - Sharing data implies sharing keys
- May interfere with standard management and recovery tasks
 - Content analysis, deduplication, indexing

Approaches



Encryption in Applications

Information is transformed by each application

- Little or no integration with other applications
- Usually, it is clear what is secure or not
 - Specific files with known file extensions

Present vulnerability windows

- Data must be decrypted to other files before being accessed

Information may be processed by different algorithms/keys

- Adapted to a specific operating system or the security level
- May complicate the data recovery processes

May difficult sharing data inside the encrypted package

- May imply extract data which is stored in a clear format

Examples:

- PGP, AxCrypt, TrueCrypt, Veracrypt, etc.
- Also: RAR, ZIP, 7Zip, LZMA...

Encryption in the File Systems

Information is transformed when is sent from memory to the filesystem

- May be broad, from the entire filesystem into the global memory cache
 - No protection in shared servers as data is available to all applications
 - Security mechanism is harder to implement in distributed environments
 - Coordination of ACLs
- May be specific to the cache of a specific process
 - Protection in the case of shared servers as data access is context-bound
 - Client API decrypts data

Examples

- EncFS, EXT4, NTFS, CFS

Encryption at the volume level

Information is transformed by the volume driver

- Transparent to applications and almost transparent to the OS
- Requires support through a specific driver
- The entire volume will be made available (partition)

Policies defined through applications or the controller

- Agnostic to the actual filesystem on top
- Protects everything, including metadata
- But it doesn't differentiate between individual users
-

Unable to solve problems related with distributed systems, but solves those related with mobile devices

- Distributed systems expose the filesystem after decryption
- Mobile devices: lost or stolen devices will keep data secure

Examples:

- PGPDisk, LUKS, BitLocker, Filevault

Encryption at the Device Level

Block Device applies security policy internally

- At boot, the device must be unlocked
- After the correct credentials are provided
- Encryption is implemented at the hardware/firmware

Advantages

- No performance loss
- Data access is not trivial as keys are internal
- May be coordinated with applications (e.g., USB devices)

Disadvantages

- After the device is unlocked, all data is made available
- Security is limited by the algorithms present
- The possible existence of backdoors is difficult to find and correct



Encryption at the Device Level

Devices have two distinct areas

- Shadow Disk: Read-Only, ~100MB with software to unlock it
- Real Disk: Read/Write. Contains user data

Two keys used

- KEK: Key Encryption Key (Authentication Key)
 - Provided by the user. Digest stored in the Shadow Disk
- MEK (or DEK): Media (Data) Encryption Key
 - Encrypted with the KEK

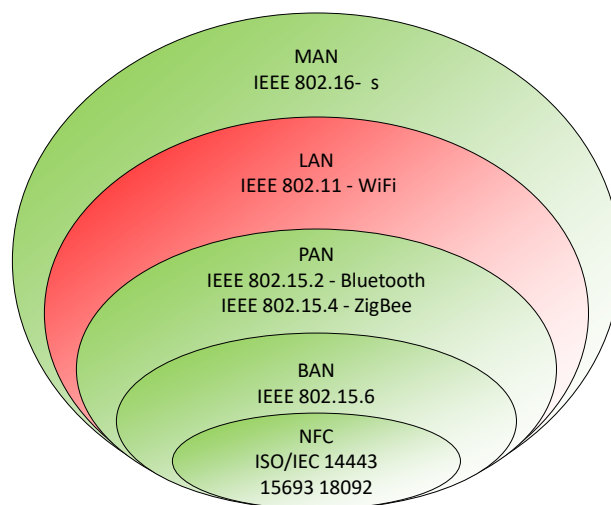
Boot process

- BIOS will access Shadow Disk and boots
- Application in Shadow Disk requests password, decrypts KEK and verifies hash(KEK)
- If it matches, MEK is decrypted, and disk geometry is updated



Security in 802.11 wireless networks

Wireless (data) communications: A glance



Wireless vs. cabled communications: Security issues

Broadcast communication

- Hard to enforce physical propagation boundaries
- Typical physical boundaries are useless to avoid:
 - Interference with communications
 - Eavesdropping of communications

Mitigation

- Reduce interference and eavesdropping capabilities
 - At the physical layer
 - At the data link layer

Reduce interference and eavesdropping capabilities: Physical layer

Prevent eavesdroppers from decoding the channel

- Channel coding needs to use some shared secret

Example: Bluetooth FHSS (Frequency Hoping Spread Spectrum)

- Carrier changes frequency in a pattern known to both transmitter and receiver
 - The data is divided into packets and transmitted over 79 hop frequencies in a pseudo random pattern
 - Only transmitters and receivers that are synchronized on the same hop frequency pattern will have access to the transmitted data
- FHSS appears as short-duration impulse noise to eavesdroppers
 - The transmitter switches hop frequencies 1,600 times per second to assure a high degree of data security

Reduce interference and eavesdropping capabilities: Physical layer

Present channel monopolization by transmitters

- Physical Medium access Policies

Examples

- Bluetooth FHSS
 - Unsynchronized transmitters seldom collide
- Wi-Fi
 - Each network is instantiated over a specific frequency
- GSM
 - Each terminal transmits over a specific mobile station

Interference is still possible from external sources or overlapping channels

Reduce interference and eavesdropping capabilities: data layer

Prevent attackers from identifying the participants in a communication

- Headers need to be encrypted, and temporary identifiers should be used

Prevent eavesdroppers from understanding data link payloads

- Frames need to be encrypted
- Usually, payloads only are encrypted

Prevent attackers from forging acceptable data link frames

- Frames need to be authenticated
 - Origin authentication
 - Freshness

IEEE 802.11: Architecture (in structured networks)

Station (STA)

- Device that can connect to a wireless network
- Has a (unique) identifier
 - Media Access Control (MAC) address
 - Today it is becoming popular its randomization (for anonymity sake)

Access Point (AP)

- Device that allows the interconnection between a wireless network and other network devices or networks

Wireless network

- Network formed by a set of STAs and AP that communicate using radio signals

IEEE 802.11: Structured network terminology

Basic Service Set (BSS)

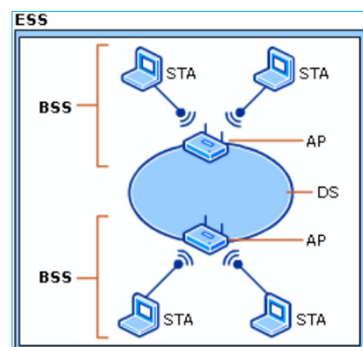
- Network formed by a set of STA associated to an AP

Extended Service Set (ESS)

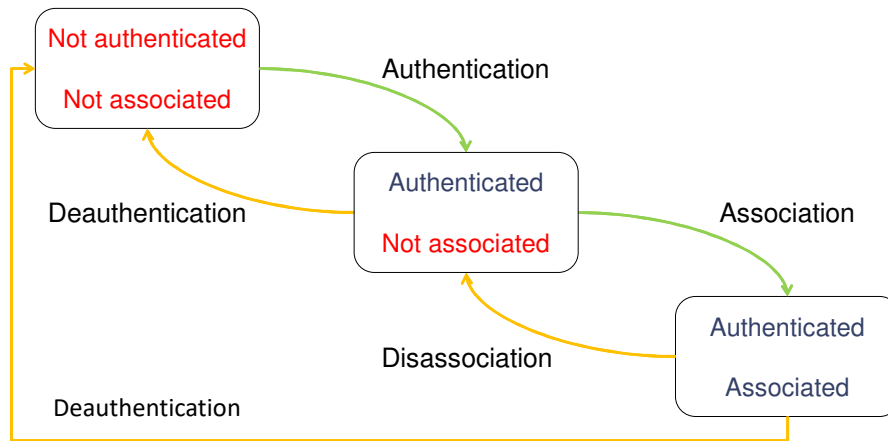
- Network formed by several BSS interconnected by a Distribution System (DS)

Service Set ID (SSID)

- Identifier of a wireless network served by a BSS or ESS
- The same infrastructure can use several SSID



IEEE 802.11: Authentication & Association state machine



© André Zuquete, João Paulo Barraca

Information and Organizational Security

10

IEEE 802.11: Frame types

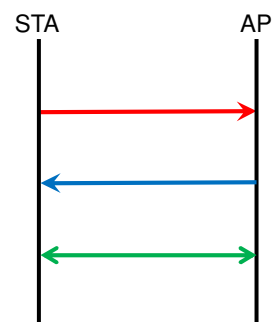
Management frames

- Beacon
- Probe Request & Response
- Authentication Request & Response
- Deauthentication
- Association Request & Response
- Reassociation Request & Response
- Disassociation

Control frames

- Request to Send (RTS)
- Clear to Send (CTS)
- Acknowledgment (ACK)

Data Frames



© André Zuquete, João Paulo Barraca

Information and Organizational Security

11

IEEE 802.11 data link security: Overview

Network Type		pre-RSN	RSN (Robust Security Network)	
Functionality		WEP	WPA	802.11i (ou WPA2)
Authentication		Unilateral (STA)	Bilateral with 802.1X (STA, AP and network)	
Key Distribution			EAP ou PSK, 4-Way Handshake	
IV Management Policy			TKIP	AES-CCMP
Data Cipher			RC4	AES-CTR
Integrity Control	Headers		Michael	AES
	Payload	CRC-32	CRC-32, Michael	CBC-MAC

Other

- SSID hiding (on beacons)
- MAC address filtering (on associations)
- (Privacy) MAC client randomization before association

IEEE 802.11: WEP (Wired Equivalent Privacy)

Optional and unilateral Authentication

- Can support multiple types simultaneously

OSA: Open System Authentication

- No authentication, just for the state transition model

SKA: Shared Key Authentication

- Challenge/response between STA and AP
- Key (password) per person (MAC address) or network
- Unilateral STA authentication
 - No AP / network authentication

Frame payload encryption

- With RC4, using 40 or 104 bit keys

Frame payload authentication with CRC-32

WEP: Lots of security problems ...

SKA is completely insecure

- An eavesdropper gets all it needs to impersonate a victim
- No need to discover the password
- Rogue APs cannot be detected

Same key for authentication and payload confidentiality

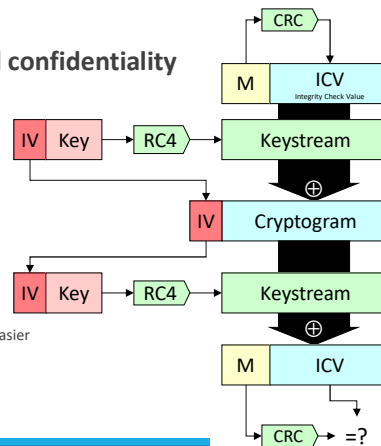
- No key distribution, keys overused

Weak integrity control

- CRC-32 is linear
- Frame deterministic modification is trivial

Mediocre IV management

- IV is too short (24 bits)
- Easy to get cryptograms produced with the same IV
- Same IV, same key \Rightarrow same keystream, cryptanalysis becomes easier
- IV is not managed at all
- Reuse is not controlled / prevented



© André Zuquete, João Paulo Barraca

Information and Organizational Security

Mitigation of WEP problems: WPA (WiFi Protected Access)

WPA uses WEP in a safe way

- A different RC4 key per frame
- RC4 weak keys are avoided
- Extra cryptographic integrity control with Michael
- IV strict sequencing for preventing frame reuse

Implemented first by device drivers

- Latter on firmware

Inline with 802.11i

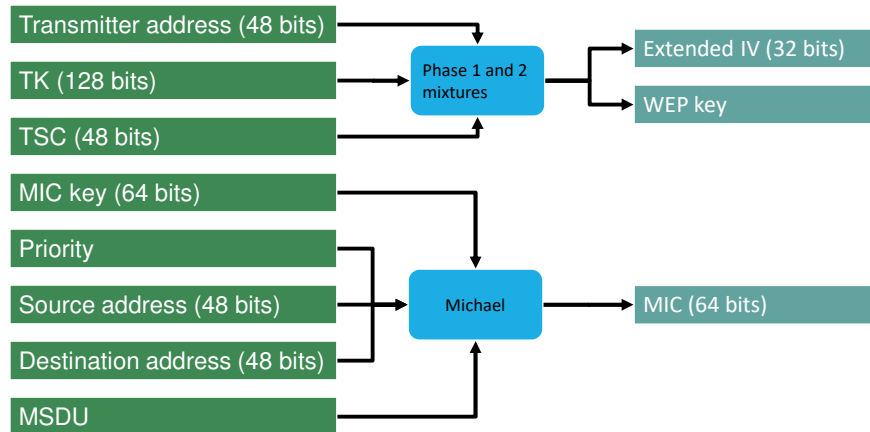
- The actual 802.11 security standard
- WPA can be used with 802.1X for strong, mutual authentication

© André Zuquete, João Paulo Barraca

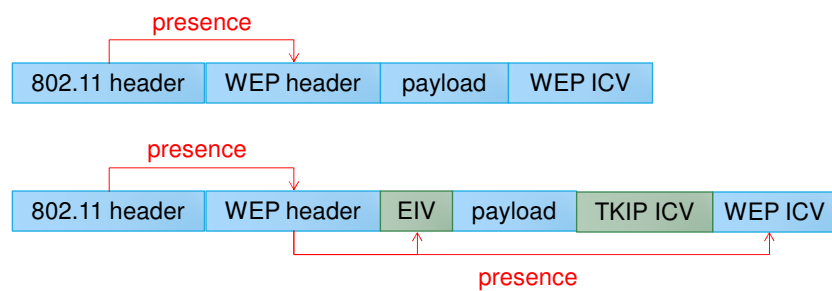
Information and Organizational Security

17

WPA: TKIP (Temporal Key Integrity Protocol)



TKIP: Frame layout



IEEE 802.1X: Port-Based Authentication

Authentication model for all IEEE 802 networks

- Layer 2 mutual authentication

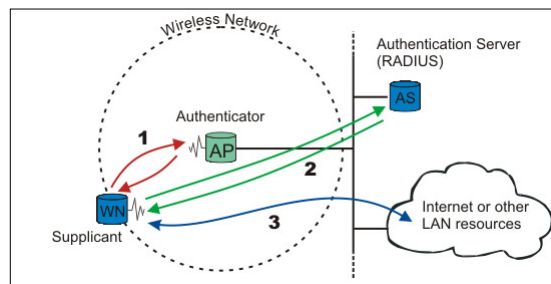
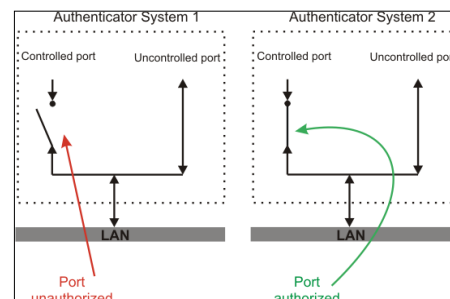
Originally conceived for large networks

- University campus, etc.
- Model was extended for wireless networks

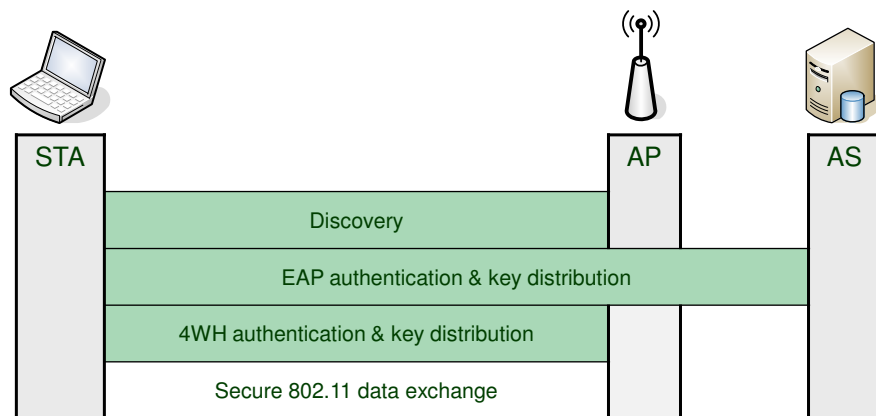
Performs key distribution

- Additional protocols focus in the remaining processes

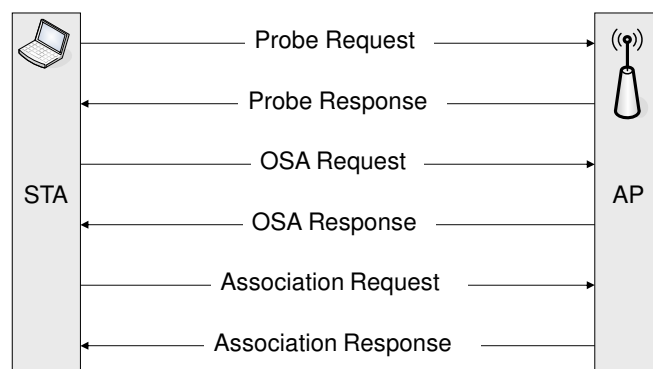
IEEE 802.1X: Architecture



IEEE 802.1X: Operational Phases



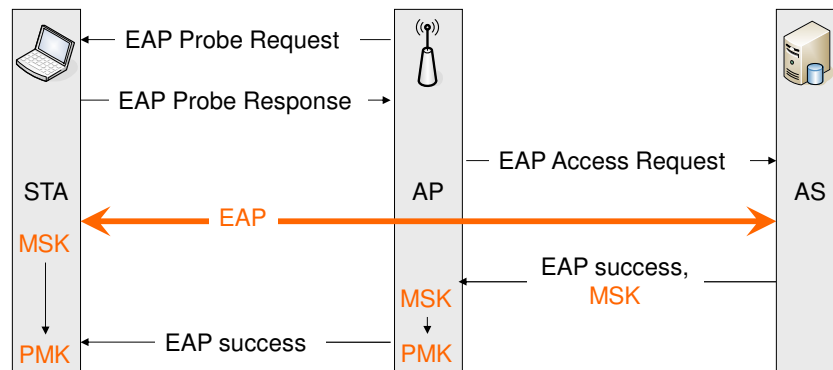
IEEE 802.1X Phase 1: Discovery (802.11 messages)



STA only got access to the AP

- 802.1X controlled port still closed

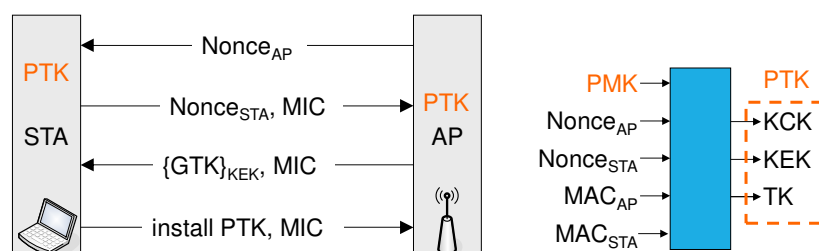
IEEE 802.1X Phase 2: Authentication (EAP Messages)



At the end of this phase AP and STA share crypto data

- **PMK** (Pairwise Master Key)
- But 802.1X controlled port still closed

IEEE 802.1X Phase 3: 4-Way Handshake (EAPoL Messages)



At the end AP and STA share new, fresh crypto data

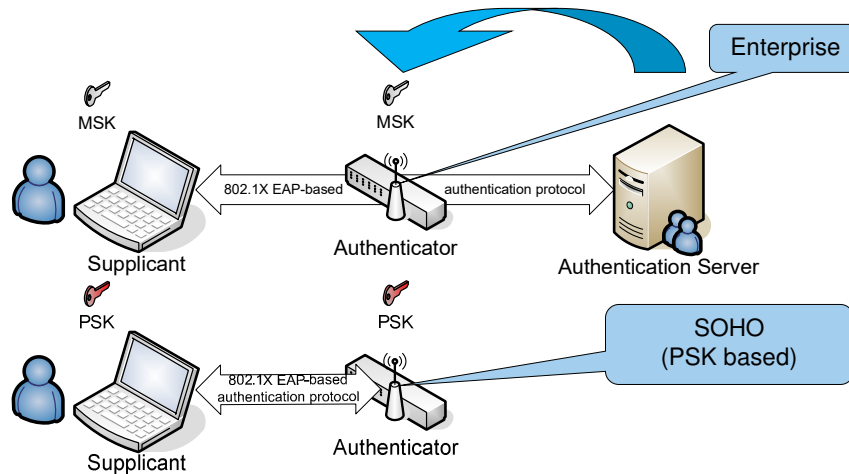
- **PTK** (Pairwise Transient Key)
- **GTK** (Group Transient Key)

Both are convinced that the peer knows **PMK and **PTK****

- Due to the use of MICs

802.1X controlled port is now open for unicast traffic

IEEE 802.1X: Architectural options

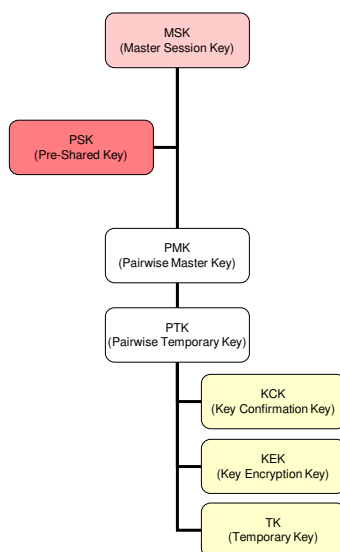


© André Zuquete, João Paulo Barraca

Information and Organizational Security

29

IEEE 802.1X: Complete key hierarchy



MSK

- Fresh outcome of an EAP protocol run
- Enterprise architecture

PSK

- Long-term AP-STA pre-shared key
- SOHO architecture

PMK

- Fresh key used for AP-STA mutual authentication and for key distribution in 4WH protocol runs

PTK

- Key used to protect AP-STA data exchanges
- KCK / KEK: 4WH protocol
- TK: 802.11 data frames

© André Zuquete, João Paulo Barraca

Information and Organizational Security

30

EAP (Extensible Authentication Protocol)

Initially conceived for PPP

- Adapted to 802.1X

AP not involved

- Relay EAP traffic
- Different EAP protocols do not imply changes in Aps

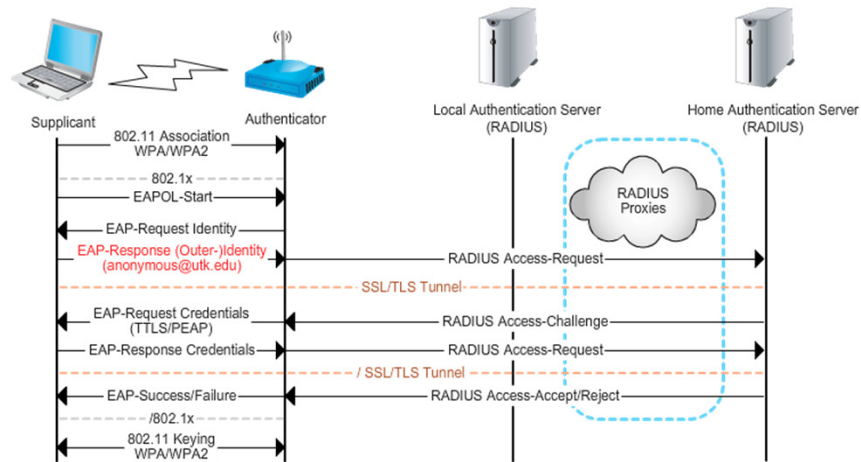
Not conceived for wireless networks

- EAP traffic not protected
- Mutual authentication not mandatory
 - An STA can be fooled by a stronger (radio level), rogue AP

Some EAP protocols for 802.1X

	LEAP	EAP-TLS	EAP-TTLS	PEAP
AS authentication	digest (challenge, password)	Public Key (certificate)		
Supplicant authentication	digest (challenge, password)	Public Key (certificate)	EAP, Public Key (certificate)	PAP, CHAP, MS-CHAP, EAP
Risks	Identity exposure Dictionary attacks Host-in-the-Middle attacks	Identity exposure		Possible identity exposure in phase 1

Eduroam: 802.1X w/ PEAP + MS-CHAPv2



Available on most University of the world

- Local Authentication Servers (using RADIUS) for roaming access

© André Zuquete, João Paulo Barraca

Information and Organizational Security

33

IEEE 802.11i (WPA2)

Defines Robust Security Networks (RSN)

- Those that support WPA and 802.11i

Uses advanced security mechanisms for frame protection

- Advanced Security Algorithm (AES) for payload encryption and frame integrity control

Uses 802.1X for network access authentication

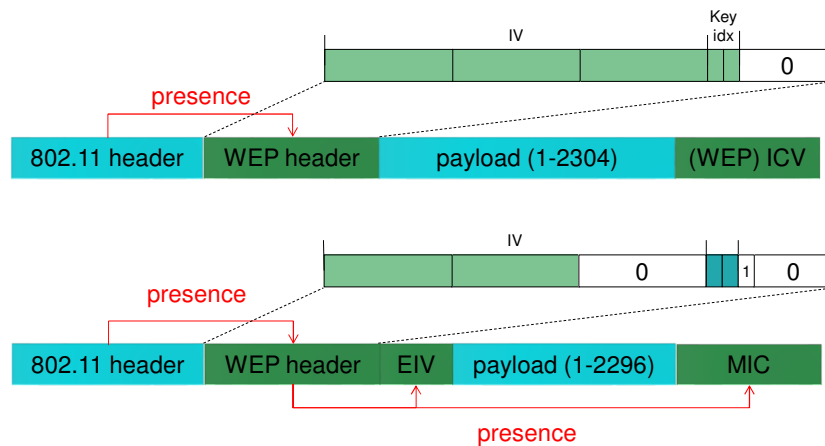
- Simplified Pre-Shared Key (PSK) mode for SOHO (Small Office, Home Office) environments
- EAP-based protocol for enterprise environments

© André Zuquete, João Paulo Barraca

Information and Organizational Security

34

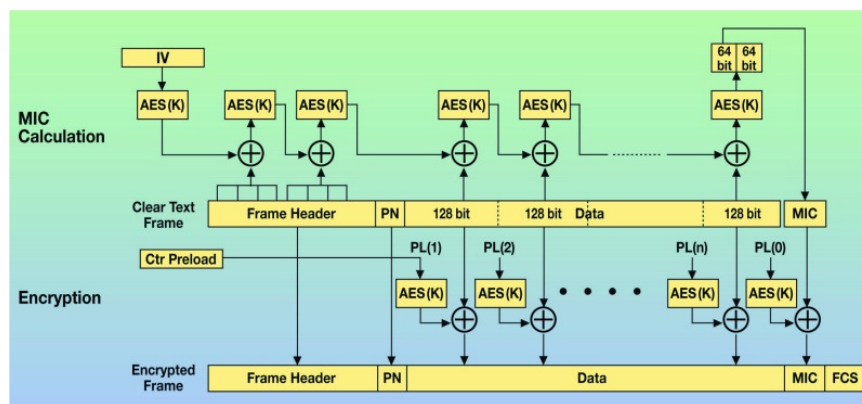
WEP vs AES-CCMP: Frame layout



WPA2 frame protection

CCMP - Counter CBC-MAC Protocol

- 128-bit keys, protection of headers, data, with cipher and authentication



<http://2014.kes.info/archiv/online/04-5-036.htm>

802.11w: Protected Management Frames

Management frames that can be used for DoS attacks are authenticated

- Deauthentication & Deassociation requests
- Other management frames unicast or broadcast by an AP

BIP (Broadcast Integrity Protocol)

- IGTK (Integrity GTK)
- For protecting part of the AP broadcast traffic

Security Association Query Request / Response

- Help to deal with desynchronization issues

IEEE 802.11 security: Are all the problems solved? No!

Dictionary attacks are still possible with PSK or EAP-based authentication

- And they will continue to be as long as (weak) passwords are chosen by people

There are still some unprotected frames

Some weaknesses at the CSMA level

- Low Congestion Window (CW) values allow attackers to get all the bandwidth